P4 Tutorial Welcome

Noa Zilberman
University of Cambridge

Goals

- **Learn P4 Language**
  - Traditional applications
  - Novel applications
- **Learn P4 software tools**
  - P4 Compiler
  - BMv2
  - P4Runtime
- **Learn about P4 hardware targets**
  - mini-workshop featuring solutions by Barefoot, Netronome, Netcope and NetFPGA.
- **Networking (the other kind)**
- **Have fun!**

Introduction to Data Plane Programming

Language Basics

Break

Software Tools & P4 Runtime

Lunch

Group Programming Exercise

Break

Architectures Mini-Workshop

Reception

Instructors

Nate Foster
Cornell

Theo Jepsen
USI Lugano

Milad Sharif
Barefoot Networks

Robert Soulě
USI Lugano

Pietro Bressana
USI Lugano

Thank you

- **Instructors & TAs**
  - Nate Foster
  - Theo Jepsen
  - Milad Sharif
  - Robert Soulě
  - Pietro Bressana
- **Mini-Workshop speakers**
  - Pavel Benacek
  - Jaco Joubert
- **You!**

P4 Language Tutorial

What is Data Plane Programming?

- **Why program the Data Plane?**

2

Software Defined Networking: Logically Centralized Control

- **Main contributions**
  - OpenFlow = standardized *model*
    - match/action abstraction
  - OpenFlow = standardized *protocol* to interact with switch
    - download flow table entries, query statistics, etc.
  - *Concept* of *logically* centralized control via a single entity ("SDN controller")
    - Simplifies control plane – e.g. compute optimal paths at one location (controller), vs. waiting for distributed routing algorithms to converge

- **Issues**
  - Data-plane protocol evolution requires changes to standards (12 → 40 OpenFlow match types)
  - Limited interoperability between vendors => southbound I/F differences handled at controller (OpenFlow / netconf / JSON / XML variants)

Control Function ("App") x m

Northbound I/F

SDN Controller x 1

Southbound I/F

Switch

x n

Status Quo: Bottom-up design

Network Demands

Switch OS

Run-time API

Driver

Fixed-function ASIC

"This is *how I know* to process packets"
(i.e. the ASIC datasheet makes the rules)

4

A Better Approach: Top-down design

Network Demands

Switch OS

Feedback

Run-time API

Driver

P4

"This is *how I want* the network to behave and how to switch packets…"
(the user / controller makes the rules)

P4 Programmable Device

5

Benefits of Data Plane Programmability

- **New Features** – Add new protocols

- **Reduce complexity** – Remove unused protocols

- **Efficient use of resources** – flexible use of tables

- **Greater visibility** – New diagnostic techniques, telemetry, etc.

- **SW style development** – rapid design cycle, fast innovation, fix data plane bugs in the field

- **You keep your own ideas**

Think programming rather than protocols…

6

Programmable Network Devices

- **PISA: Flexible Match+Action ASICs**
  - ◦ Intel Flexpipe, Cisco Doppler, Cavium (Xpliant), Barefoot Tofino, …
- **NPU**
  - ◦ EZchip, Netronome, …
- **CPU**
  - ◦ Open Vswitch, eBPF, DPDK, VPP…
- **FPGA**
  - ◦ Xilinx, Altera, …

These devices let us tell them how to process packets.

What can you do with P4?

- **Layer 4 Load Balancer – SilkRoad[1]**

- **Low Latency Congestion Control – NDP[2]**

- **In-band Network Telemetry – INT[3]**

- **Fast In-Network cache for key-value stores – NetCache[4]**

- **Consensus at network speed – NetPaxos[5]**

- **Aggregation for MapReduce Applications [6]**

- **… and much more**

[1] Miao, Rui, et al. "SilkRoad: Making Stateful Layer-4 Load Balancing Fast and Cheap Using Switching ASICs." SIGCOMM, 2017.
[2] Handley, Mark, et al. "Re-architecting datacenter networks and stacks for low latency and high performance." SIGCOMM, 2017.
[4] Kim, Changhoon, et al. "In-band network telemetry via programmable dataplanes." SIGCOMM. 2015.
[3] Xin Jin et al. "NetCache: Balancing Key-Value Stores with Fast In-Network Caching." To appear at SOSP 2017
[5] Dang, Huynh Tu, et al. "NetPaxos: Consensus at network speed." SIGCOMM, 2015.
[6] Sapio, Amedeo, et al. "In-Network Computation is a Dumb Idea Whose Time Has Come." *Hot Topics in Networks*. ACM, 2017.

Brief History and Trivia

- **May 2013: Initial idea and the name "P4"**
- **July 2014: First paper (SIGCOMM CCR)**
- **Aug 2014: First P4$_{14}$ Draft Specification (v0.9.8)**
- **Sep 2014: P4$_{14}$ Specification released (v1.0.0)**
- **Jan 2015: P4$_{14}$ v1.0.1**
- **Mar 2015: P4$_{14}$ v1.0.2**
- **Nov 2016: P4$_{14}$ v1.0.3**
- **May 2017: P4$_{14}$ v1.0.4**

- **Apr 2016: P4$_{16}$ – first commits**
- **Dec 2016: First P4$_{16}$ Draft Specification**
- **May 2017: P4$_{16}$ Specification released**

P4_16 Data Plane Model

PISA: Protocol-Independent Switch Architecture

Programmer declares the headers that should be recognized and their order in the packet

Programmer defines the tables and the exact processing algorithm

Programmer declares how the output packet will look on the wire

Programmable Parser

Programmable Match-Action Pipeline
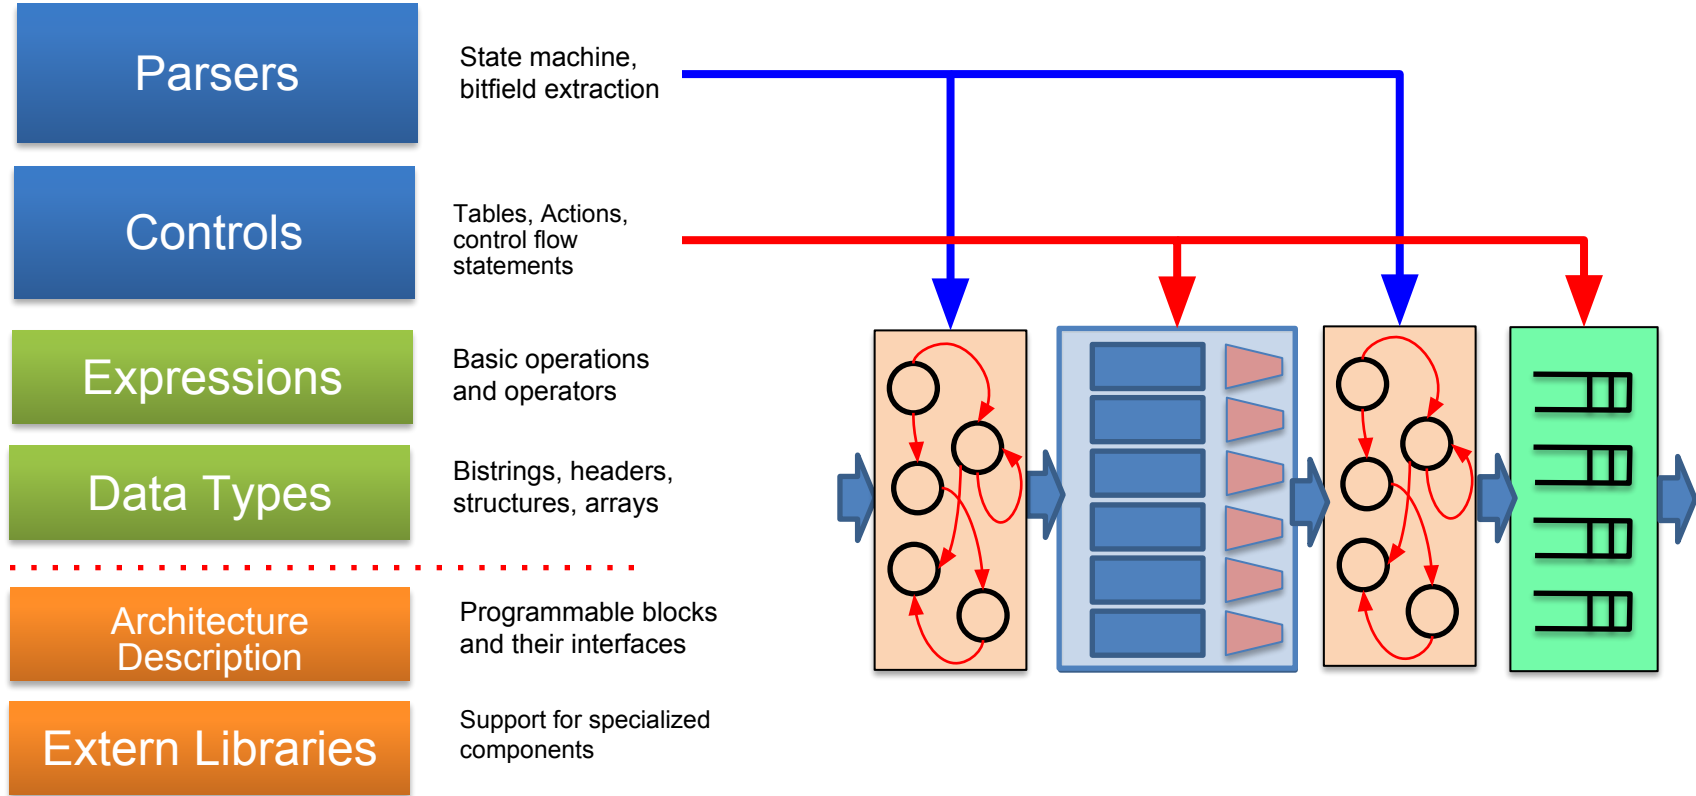
Programmable Deparser
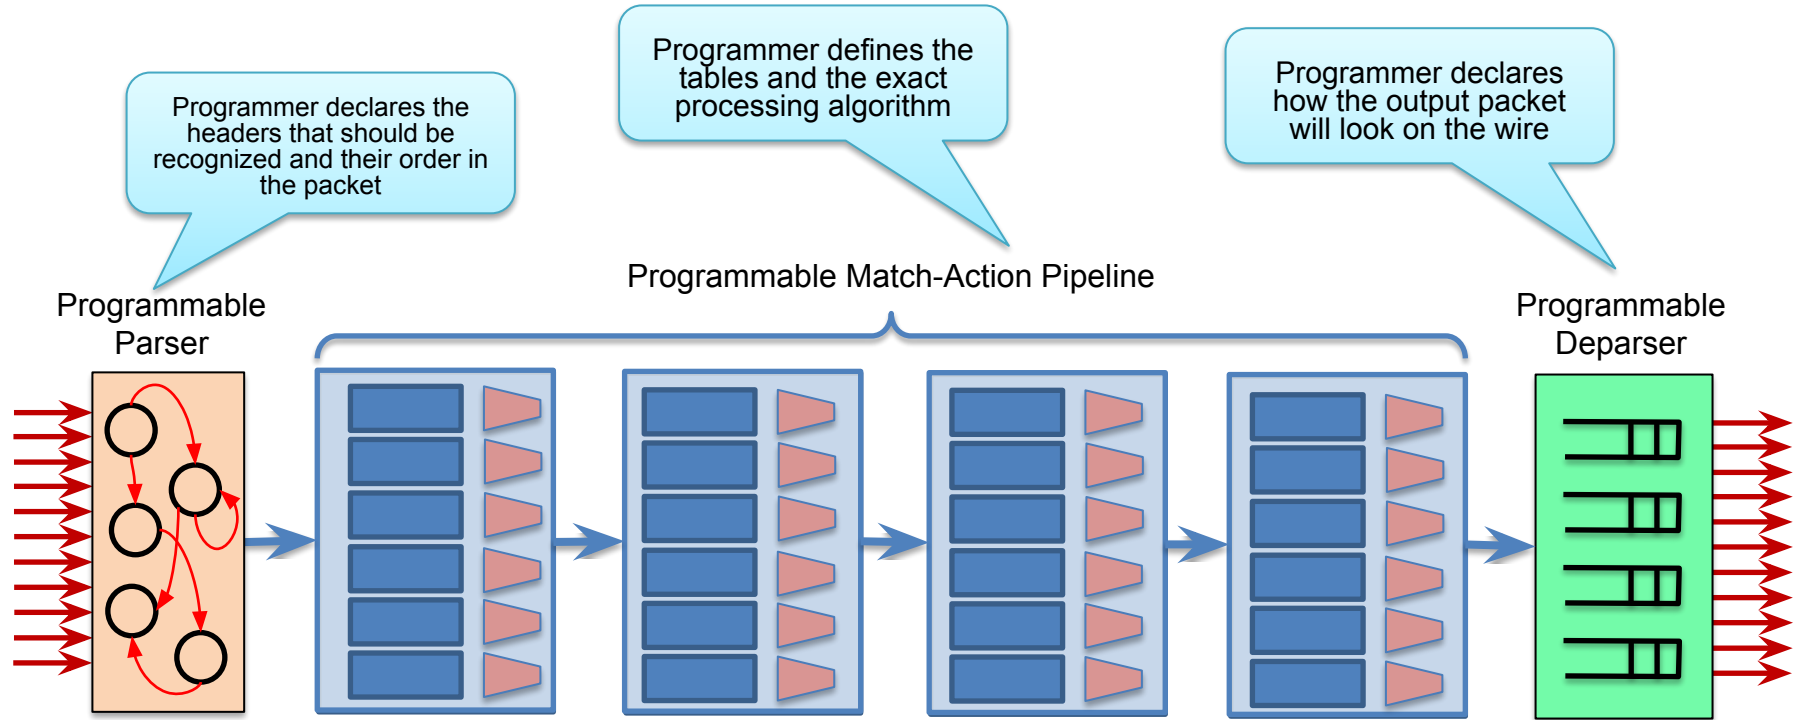
PISA in Action

- **Packet is parsed into individual headers (parsed representation)**
- **Headers and intermediate results can be used for matching and actions**
- **Headers can be modified, added or removed**
- **Packet is deparsed (serialized)**

Programmable Parser

Programmable Match-Action Pipeline

Programmable Deparser

P4₁₆ Language Elements

Element	Description
Parsers	State machine, bitfield extraction
Controls	Tables, Actions, control flow statements
Expressions	Basic operations and operators
Data Types	Bistrings, headers, structures, arrays
Architecture Description	Programmable blocks and their interfaces
Extern Libraries	Support for specialized components

P4_16 Approach

| Term	Explanation
P4 Target	An embodiment of a specific hardware implementation
P4 Architecture	Provides an interface to program a target via some set of P4-programmable components, externs, fixed components

Example Architectures and Targets

V1Model

SimpleSumeSwitch

Portable Switch Architecture (PSA)

Anything

15

# Programming a P4 Target

Lab 1: Basics

Before we start...

- **Install VM image (Look for instructor with USB sticks)**
- **Please make sure that your VM is up to date**
  - ◦ `$ cd ~/tutorials && git pull`
- **We'll be using several software tools pre-installed on the VM**
  - ◦ Bmv2: a P4 software switch
  - ◦ p4c: the reference P4 compiler
  - ◦ Mininet: a lightweight network emulation environment
- **Each directory contains a few scripts**
  - ◦ `$ make` : compiles P4 program, execute on Bmv2 in Mininet, populate tables
  - ◦ `*.py`: send and receive test packets in Mininet
- **Exercises**
  - ◦ Each example comes with an incomplete implementation; your job is to finish it!
  - ◦ Look for "TODOs" (or peek at the P4 code in `solution/` if you must)

18

V1Model Architecture

- **Implemented on top of Bmv2's** `simple_switch` **target**

V1Model Standard Metadata

```
struct standard_metadata_t {
    bit<9>  ingress_port;
    bit<9>  egress_spec;
    bit<9>  egress_port;
    bit<32> clone_spec;
    bit<32> instance_type;
    bit<1>  drop;
    bit<16> recirculate_port;
    bit<32> packet_length;
    bit<32> enq_timestamp;
    bit<19> enq_qdepth;
    bit<32> deq_timedelta;
    bit<19> deq_qdepth;
    bit<48> ingress_global_timestamp;
    bit<32> lf_field_list;
    bit<16> mcast_grp;
    bit<1>  resubmit_flag;
    bit<16> egress_rid;
    bit<1>  checksum_error;
}
```

- **ingress_port** - the port on which the packet arrived

- **egress_spec** - the port to which the packet should be sent to

- **egress_port** - the port on which the packet is departing from (read only in egress pipeline)

P4₁₆ Program Template (V1Model)

```
#include <core.p4>
#include <v1model.p4>
/* HEADERS */
struct metadata { ... }
struct headers {
  ethernet_t   ethernet;
  ipv4_t       ipv4;
}
/* PARSER */
parser MyParser(packet_in packet,
        out headers hdr,
        inout metadata meta,
        inout standard_metadata_t smeta) {
  ...
}
/* CHECKSUM VERIFICATION */
control MyVerifyChecksum(in headers hdr,
                  inout metadata meta) {
  ...
}
/* INGRESS PROCESSING */
control MyIngress(inout headers hdr,
            inout metadata meta,
            inout standard_metadata_t std_meta) {
  ...
}
```

```
/* EGRESS PROCESSING */
control MyEgress(inout headers hdr,
            inout metadata meta,
            inout standard_metadata_t std_meta) {
  ...
}
/* CHECKSUM UPDATE */
control MyComputeChecksum(inout headers hdr,
                  inout metadata meta) {
  ...
}
/* DEPARSER */
control MyDeparser(inout headers hdr,
            inout metadata meta) {
  ...
}
/* SWITCH */
V1Switch(
  MyParser(),
  MyVerifyChecksum(),
  MyIngress(),
  MyEgress(),
  MyComputeChecksum(),
  MyDeparser()
) main;
```

P4₁₆ Hello World (V1Model)

```
#include <core.p4>
#include <v1model.p4>
struct metadata {}
struct headers {}

parser MyParser(packet_in packet,
    out headers hdr,
    inout metadata meta,
    inout standard_metadata_t standard_metadata) {

     state start { transition accept; }
}

control MyVerifyChecksum(inout headers hdr, inout metadata
meta) {   apply {  }   }

control MyIngress(inout headers hdr,
    inout metadata meta,
    inout standard_metadata_t standard_metadata) {
apply {
    if (standard_metadata.ingress_port == 1) {
        standard_metadata.egress_spec = 2;
    } else if (standard_metadata.ingress_port == 2) {
        standard_metadata.egress_spec = 1;
    }
    }
}
```

```
control MyEgress(inout headers hdr,
    inout metadata meta,
    inout standard_metadata_t standard_metadata) {
    apply {  }
}

control MyComputeChecksum(inout headers hdr, inout metadata
meta) {
    apply { }
}

control MyDeparser(packet_out packet, in headers hdr) {
    apply { }
}

V1Switch(
    MyParser(),
    MyVerifyChecksum(),
    MyIngress(),
    MyEgress(),
    MyComputeChecksum(),
    MyDeparser()
) main;
```

P4₁₆ Hello World (V1Model)

```
#include <core.p4>
#include <v1model.p4>
struct metadata {}
struct headers {}

parser MyParser(packet_in packet, out headers hdr,
    inout metadata meta,
    inout standard_metadata_t standard_metadata) {
     state start { transition accept; }
}

control MyIngress(inout headers hdr, inout metadata meta,
    inout standard_metadata_t standard_metadata) {
     action set_egress_spec(bit<9> port) {
         standard_metadata.egress_spec = port;
     }
     table forward {
         key = { standard_metadata.ingress_port: exact; }
         actions = {
             set_egress_spec;
             NoAction;
         }
         size = 1024;
         default_action = NoAction();
     }
     apply {   forward.apply();   }
}
```
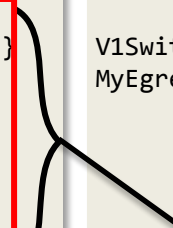
```
control MyEgress(inout headers hdr,
    inout metadata meta,
    inout standard_metadata_t standard_metadata) {
    apply {   }
}

control MyVerifyChecksum(inout headers hdr, inout metadata
meta) {   apply { }   }

control MyComputeChecksum(inout headers hdr, inout metadata
meta) {   apply { }   }

control MyDeparser(packet_out packet, in headers hdr) {
    apply { }
}

V1Switch( MyParser(), MyVerifyChecksum(), MyIngress(),
MyEgress(), MyComputeChecksum(), MyDeparser() ) main;
```

| Key	Action Name	Action Data
1	set_egress_spec	2
2	set_egress_spec	1

Running Example: Basic Forwarding

- **We'll use a simple application as a running example—a basic router—to illustrate the main features of P4$_{16}$**

- **Basic router functionality:**
  - Parse Ethernet and IPv4 headers from packet
  - Find destination in IPv4 routing table
  - Update source / destination MAC addresses
  - Decrement time-to-live (TTL) field
  - Set the egress port
  - Deparse headers back into a packet

- **We've written some starter code for you (`basic.p4`) and implemented a static control plane**

Basic Forwarding: Topology

h1
(10.0.1.1)

s1

h2
(10.0.2.2)

s2

s3

h3
(10.0.3.3)

P4₁₆ Types (Basic and Header Types)

```
typedef bit<48> macAddr_t;
typedef bit<32> ip4Addr_t;
header ethernet_t {
  macAddr_t dstAddr;
  macAddr_t srcAddr;
  bit<16>   etherType;
}
header ipv4_t {
  bit<4>     version;
  bit<4>     ihl;
  bit<8>     diffserv;
  bit<16>    totalLen;
  bit<16>    identification;
  bit<3>     flags;
  bit<13>    fragOffset;
  bit<8>     ttl;
  bit<8>     protocol;
  bit<16>    hdrChecksum;
  ip4Addr_t srcAddr;
  ip4Addr_t dstAddr;
}
```

Basic Types
- **bit<n>**: Unsigned integer (bitstring) of size n
- **bit** is the same as **bit<1>**
- **int<n>**: Signed integer of size n (>=2)
- **varbit<n>**: Variable-length bitstring

Header Types: Ordered collection of members
- Can contain **bit<n>**, **int<n>**, and **varbit<n>**
- Byte-aligned
- Can be valid or invalid
- Provides several operations to test and set validity bit: **isValid()**, **setValid()**, and **setInvalid()**

Typedef: Alternative name for a type

P4₁₆ Types (Other Types)

```
/* Architecture */
struct standard_metadata_t {
  bit<9>  ingress_port;
  bit<9>  egress_spec;
  bit<9>  egress_port;
  bit<32> clone_spec;
  bit<32> instance_type;
  bit<1>  drop;
  bit<16> recirculate_port;
  bit<32> packet_length;
  ...
}

/* User program */
struct metadata {
  ...
}
struct headers {
  ethernet_t   ethernet;
  ipv4_t       ipv4;
}
```

Other useful types

- **Struct**: Unordered collection of members (with no alignment restrictions)

- **Header Stack:** array of headers

- **Header Union:** one of several headers

P4₁₆ Parsers

- **Parsers are functions that map packets into headers and metadata, written in a state machine style**
- **Every parser has three predefined states**
  - ◦ start
  - ◦ accept
  - ◦ reject
- **Other states may be defined by the programmer**
- **In each state, execute zero or more statements, and then transition to another state (loops are OK)**

Parsers (V1Model)

```
/* From core.p4 */
extern packet_in {
  void extract<T>(out T hdr);
  void extract<T>(out T variableSizeHeader,
           in bit<32> variableFieldSizeInBits);
  T lookahead<T>();
  void advance(in bit<32> sizeInBits);
  bit<32> length();
}
/* User Program */
parser MyParser(packet_in packet,
         out headers hdr,
         inout metadata meta,
         inout standard_metadata_t std_meta) {

  state start {
    packet.extract(hdr.ethernet);
    transition accept;
  }

}
```

MyParser

packet_in   hdr

meta   meta

The platform Initializes
User Metadata to 0

standard_meta

Select Statement

```
state start {
  transition parse_ethernet;
}

state parse_ethernet {
  packet.extract(hdr.ethernet);
  transition select(hdr.ethernet.etherType) {
    0x800: parse_ipv4;
    default: accept;
  }
}
```

P4$_{16}$ has a `select` statement that can be used to branch in a parser

Similar to `case` statements in C or Java, but without "fall-through behavior"—i.e., `break` statements are not needed

In parsers it is often necessary to branch based on some of the bits just parsed

For example, etherType determines the format of the rest of the packet

Match patterns can either be literals or simple computations such as masks

Coding Break

31

P4₁₆ Controls

- **Similar to C functions (without loops)**

- **Can declare variables, create tables, instantiate externs, etc.**

- **Functionality specified by code in `apply` statement**

- **Represent all kinds of processing that are expressible as DAG:**
  - Match-Action Pipelines
  - Deparsers
  - Additional forms of packet processing (updating checksums)

- **Interfaces with other blocks are governed by user- and architecture-specified types (typically headers and metadata)**

Example: Reflector (V1Model)

```
control MyIngress(inout headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t std_meta) {
  /* Declarations region */
  bit<48> tmp;

  apply {
    /* Control Flow */
    tmp = hdr.ethernet.dstAddr;
    hdr.ethernet.dstAddr = hdr.ethernet.srcAddr;
    hdr.ethernet.srcAddr = tmp;
    std_meta.egress_spec = std_meta.ingress_port;
  }
}
```

Desired Behavior:

- **Swap source and destination MAC addresses**

- **Bounce the packet back out on the physical port that it came into the switch on**

Example: Simple Actions

```
control MyIngress(inout headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t std_meta) {

  action swap_mac(inout bit<48> src,
                  inout bit<48> dst) {
    bit<48> tmp = src;
    src = dst;
    dst = tmp;
  }

  apply {
    swap_mac(hdr.ethernet.srcAddr,
             hdr.ethernet.dstAddr);
    std_meta.egress_spec = std_meta.ingress_port;
  }
}
```

- **Very similar to C functions**
- **Can be declared inside a control or globally**
- **Parameters have type and direction**
- **Variables can be instantiated inside**
- **Many standard arithmetic and logical operations are supported**
  - +, -, *
  - ~, &, |, ^, >>, <<
  - ==, !=, >, >=, <, <=
  - No division/modulo
- **Non-standard operations:**
  - Bit-slicing: [m:l] (works as l-value too)
 - Bit Concatenation: ++

P4$_{16}$ Tables

- **The fundamental unit of a Match-Action Pipeline**
  - Specifies what data to match on and match kind
  - Specifies a list of *possible* actions
  - Optionally specifies a number of table **properties**
    - Size
    - Default action
    - Static entries
    - etc.
- **Each table contains one or more entries (rules)**
- **An entry contains:**
  - A specific key to match on
  - A **single** action that is executed when a packet matches the entry
  - Action data (possibly empty)

Tables: Match-Action Processing

Control Plane

Headers and Metadata

Lookup Key

Directional (DataPlane) Parameters

Hit

Key	Action ID	Action Data

Default Action ID	Default Action Data

Hit/Miss Selector

Action ID

Data

Action Code

Directionless (Action Data) Parameters

Action Execution

Headers and Metadata (Input)

Headers and Metadata (Output)

Example: IPv4_LPM Table

10.0.1.1   10.0.1.2

Key	Action	Action Data
10.0.1.1/32	ipv4_forward	dstAddr=00:00:00:00:01:01 port=1
10.0.1.2/32	drop	
*`	NoAction	

- **Data Plane (P4) Program**
  - Defines the format of the table
    - Key Fields
    - Actions
    - Action Data
  - Performs the lookup
  - Executes the chosen action
- **Control Plane (IP stack, Routing protocols)**
  - Populates table entries with specific information
    - Based on the configuration
    - Based on automatic discovery
    - Based on protocol calculations

37

IPv4_LPM Table

```
table ipv4_lpm {
    key = {
        hdr.ipv4.dstAddr: lpm;
    }
    actions = {
        ipv4_forward;
        drop;
        NoAction;
    }
    size = 1024;
    default_action = NoAction();
}
```

Match Kinds

```
/* core.p4 */
match_kind {
    exact,
    ternary,
    lpm
}

/* v1model.p4 */
match_kind {
    range,
    selector
}

/* Some other architecture */
match_kind {
    regexp,
    fuzzy
}
```

- **The type `match_kind` is special in P4**

- **The standard library (`core.p4`) defines three standard match kinds**
  - Exact match
  - Ternary match
  - LPM match

- **The architecture (`v1model.p4`) defines two additional match kinds:**
  - range
  - selector

- **Other architectures may define (and provide implementation for) additional match kinds**

Defining Actions for L3 forwarding

```
/* core.p4 */
action NoAction() {
}

/* basic.p4 */
action drop() {
  mark_to_drop();
}

/* basic.p4 */
action ipv4_forward(macAddr_t dstAddr,
                    bit<9> port) {
  ...
}
```

- **Actions can have two different types of parameters**
  - ◦ Directional (from the Data Plane)
  - ◦ Directionless (from the Control Plane)
- **Actions that are called directly:**
  - ◦ Only use directional parameters
- **Actions used in tables:**
  - ◦ Typically use directionless parameters
  - ◦ May sometimes use directional parameters too

Directional (DataPlane) Parameters

Action Code

Directionless (Action Data) Parameters

Action Execution

Applying Tables in Controls

```
control MyIngress(inout headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t standard_metadata) {
  table ipv4_lpm {
    ...
  }
  apply {
    ...
    ipv4_lpm.apply();
    ...
  }
}
```

P4₁₆ Deparsing

```
/* From core.p4 */
extern packet_out {
  void emit<T>(in T hdr);
}

/* User Program */
control DeparserImpl(packet_out packet,
                     in headers hdr) {
  apply {
    ...
    packet.emit(hdr.ethernet);
    ...
  }
}
```
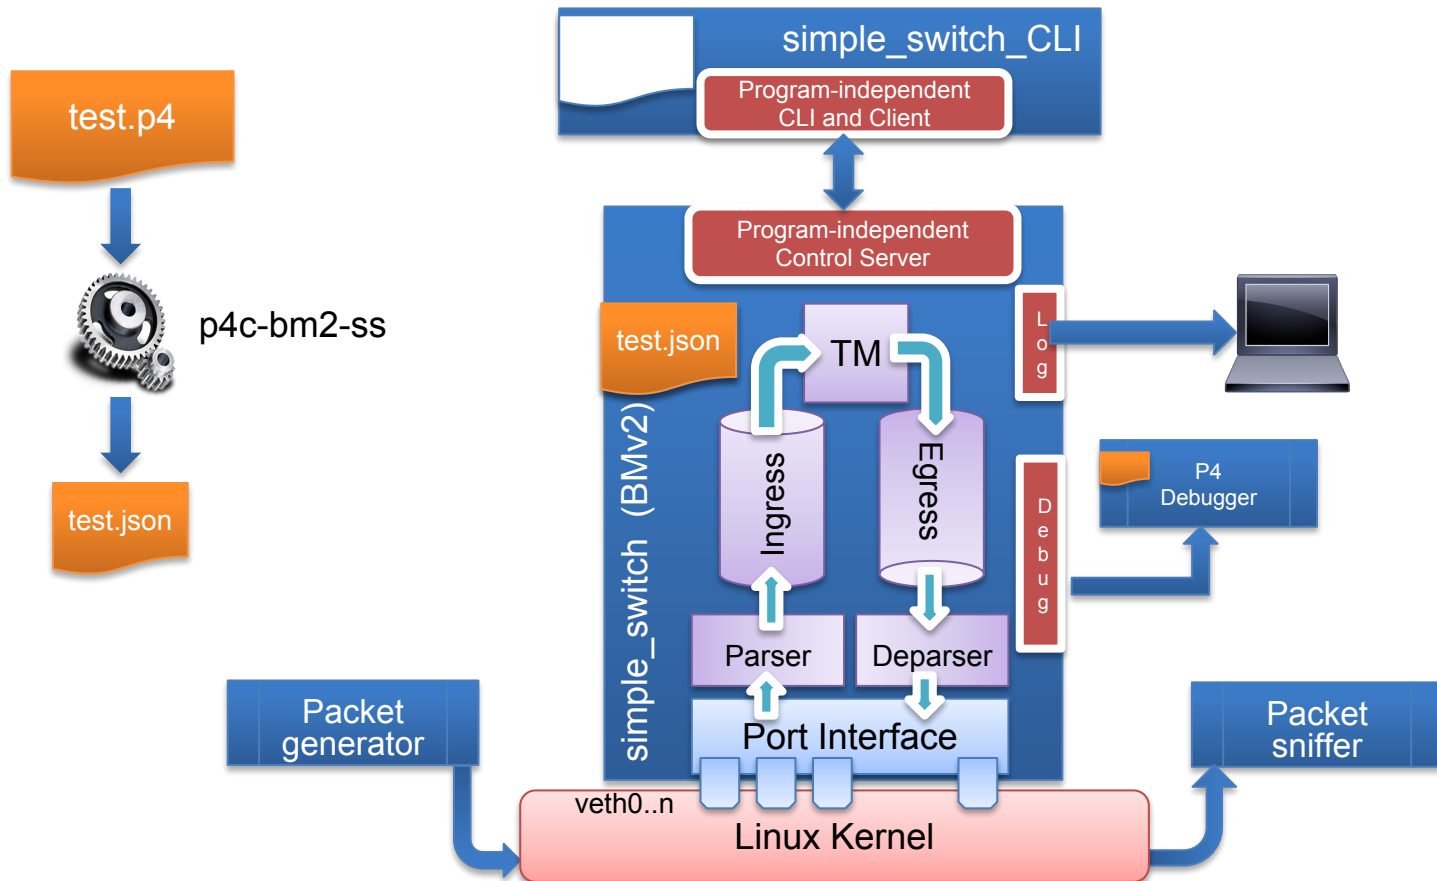
- **Assembles the headers back into a well-formed packet**

- **Expressed as a control function**
  ◦ No need for another construct!

- **packet_out extern is defined in core.p4:** emit(hdr): serializes header if it is valid

- **Advantages:**
  - Makes deparsing explicit...
   ...but decouples from parsing

Coding Break

# Makefile: under the hood

Step 1: P4 Program Compilation

test.p4

p4c-bm2-ss

test.json

```
$ p4c-bm2-ss -o test.json test.p4
```

Step 2: Preparing veth Interfaces

test.p4

test.json

```
$ sudo ~/p4lang/tutorials/examples/veth_setup.sh

# ip link add name veth0 type veth peer name veth1
# for iface in "veth0 veth1"; do
    ip link set dev ${iface} up
    sysctl net.ipv6.conf.${iface}.disable_ipv6=1
    TOE_OPTIONS="rx tx sg tso ufo gso gro lro rxvlan txvlan rxhash"
    for TOE_OPTION in $TOE_OPTIONS; do
        /sbin/ethtool --offload $intf "$TOE_OPTION"
    done
  done
```
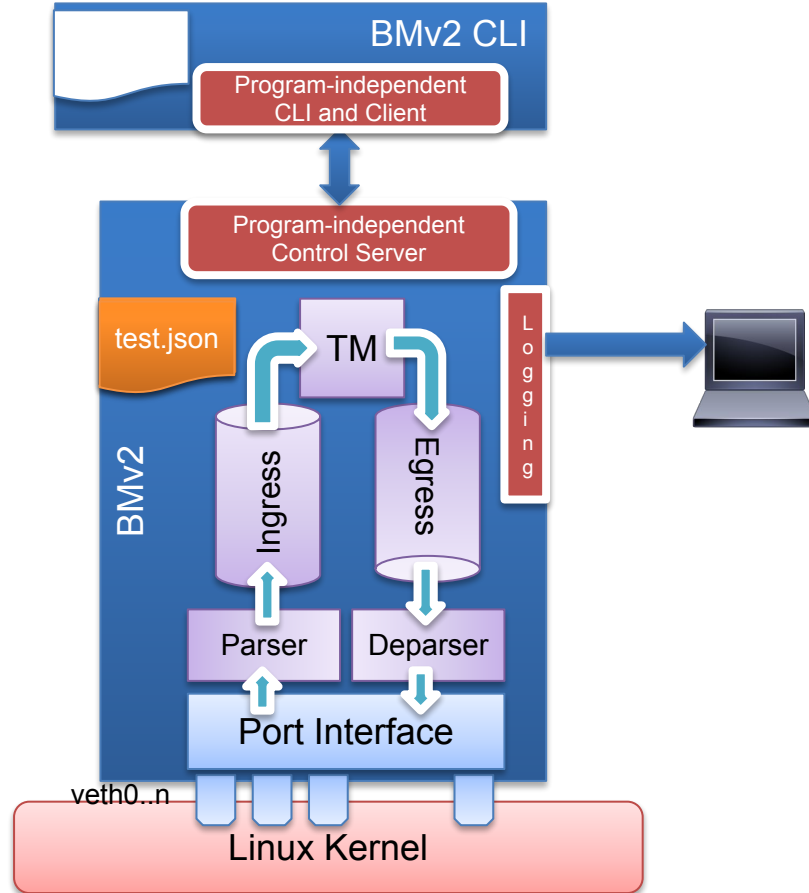
veth

0 2 4 $2n$

Linux
Kernel

1 3 5 $2n+1$

veth

Step 3: Starting the model

```
$ sudo simple_switch --log-console --dump-packet-data 64 \
    -i 0@veth0 -i 1@veth2 … [--pcap]                       \
    test.json
```
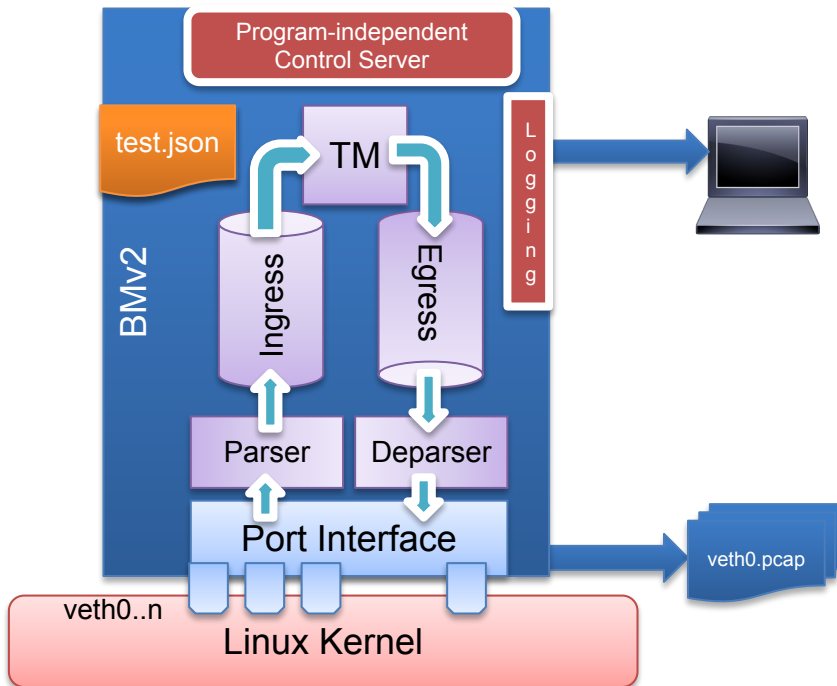
Step 4: Starting the CLI


```
$ simple_switch_CLI
```

Working with Tables in simple_switch_CLI

```
RuntimeCmd: show_tables
m_filter                        [meta.meter_tag(exact, 32)]
m_table                         [ethernet.srcAddr(ternary, 48)]

RuntimeCmd: table_info m_table
m_table                         [ethernet.srcAddr(ternary, 48)]
********************************************************************************
_nop
[]m_action                      [meter_idx(32)]

RuntimeCmd: table_dump m_table
m_table:
0: aaaaaaaaaaaa &&& ffffffffffff => m_action - 0,
SUCCESS

RuntimeCmd: table_add m_table m_action 01:00:00:00:00:00&&&01:00:00:00:00:00 => 1 0
Adding entry to ternary match table m_table
match key:            TERNARY-01:00:00:00:00:00 &&& 01:00:00:00:00:00
action:               m_action
runtime data:         00:00:00:05
SUCCESS
entry has been added with handle 1

RuntimeCmd: table_delete 1
```

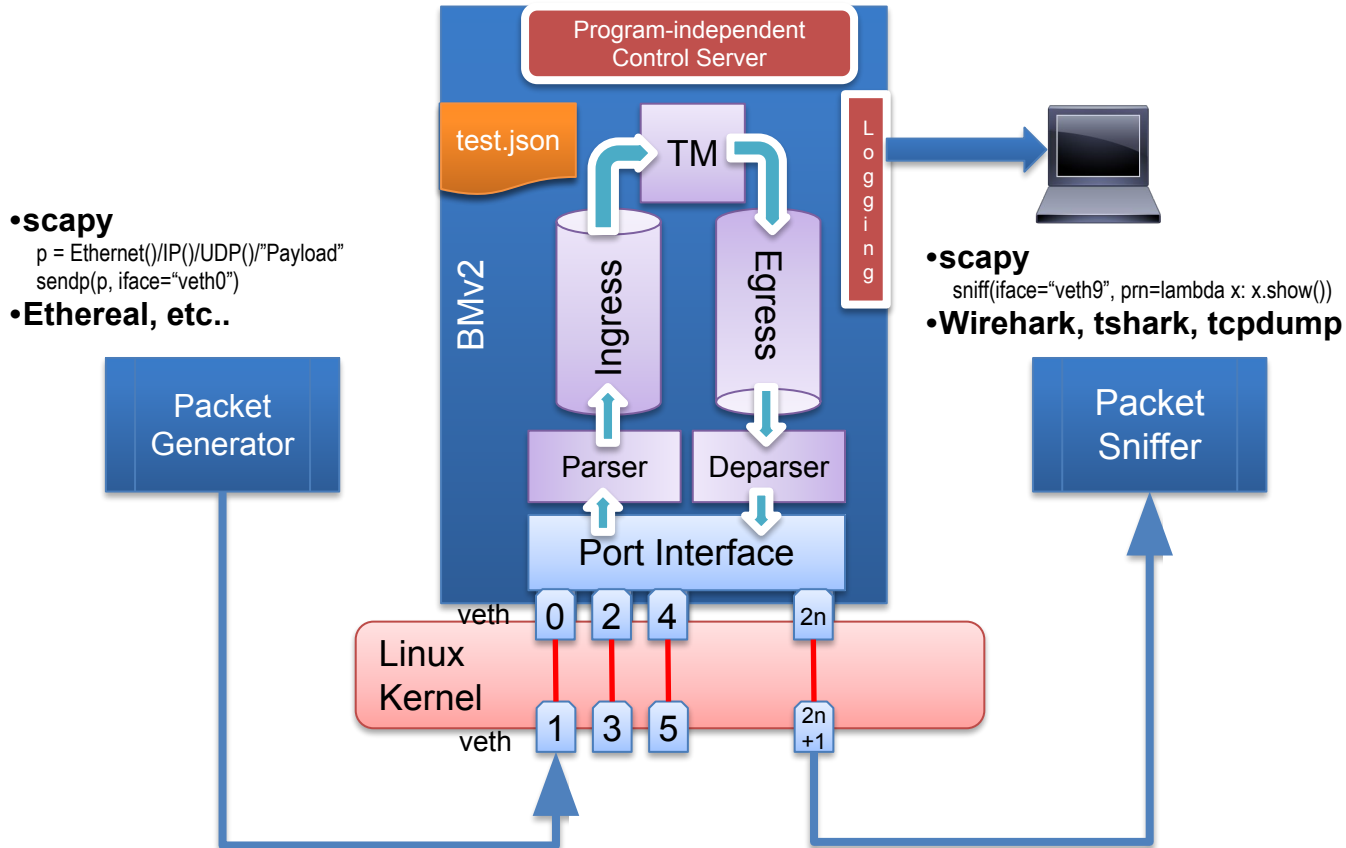Value and mask for ternary matching. No spaces around "&&&"

Entry priority

"=>" separates the key from the action data

All subsequent operations use the entry handle

Step 5: Sending and Receiving Packets

- **scapy**
  p = Ethernet()/IP()/UDP()/"Payload"
  sendp(p, iface="veth0")
- **Ethereal, etc..**

Program-independent Control Server

test.json

TM

Logging

Ingress

Egress

BMv2

Parser

Deparser

Port Interface

- **scapy**
  sniff(iface="veth9", prn=lambda x: x.show())
- **Wirehark, tshark, tcpdump**

Packet Generator

Packet Sniffer

veth 0 2 4 2n

Linux Kernel

veth 1 3 5 2n+1

50

Basic Tunneling

- Add support for basic tunneling to the basic IP router

- Define a new header type (`myTunnel`) to encapsulate the IP packet

- `myTunnel` header includes:
  - `proto_id` : type of packet being encapsulated
  - `dst_id` : ID of destination host

- Modify the switch to perform routing using the `myTunnel` header

Basic Tunneling TODO List

- **Define `myTunnel_t` header type and add to `headers` struct**

- **Update parser**

- **Define `myTunnel_forward` action**

- **Define `myTunnel_exact` table**

- **Update table application logic in `MyIngress` apply statement**

- **Update deparser**

- **Adding forwarding rules**

Basic Forwarding: Topology

h1
(10.0.1.1)
(dst_id: 1)

s1

1

2

3

h2
(10.0.2.2)
(dst_id: 2)

s2

1

2

3

s3

2

3

1

h3
(10.0.3.3) (dst_id: 3)

Coding Break

FAQs

- **Can I apply a table multiple times in my P4 Program?**

  - No (except via resubmit / recirculate)

- **Can I modify table entries from my P4 Program?**

  - No (except for direct counters)

- **What happens upon reaching the `reject` state of the parser?**

  - Architecture dependent

- **How much of the packet can I parse?**

  - Architecture dependent

Fin!

Debugging

```
control MyIngress(...) {
  table debug {
    key = {
      std_meta.egress_spec : exact;
    }
    actions = { }
  }
  apply {
    ...
    debug.apply();
  }
}
```

- **Bmv2 maintains logs that keep track of how packets are processed in detail**
  - `/tmp/p4s.s1.log`
  - `/tmp/p4s.s2.log`
  - `/tmp/p4s.s3.log`

- **Can manually add information to the logs by using a dummy debug table that reads headers and metadata of interest**

- `[15:16:48.145] [bmv2] [D] [thread 4090] [96.0] [cxt 0] Looking up key:`
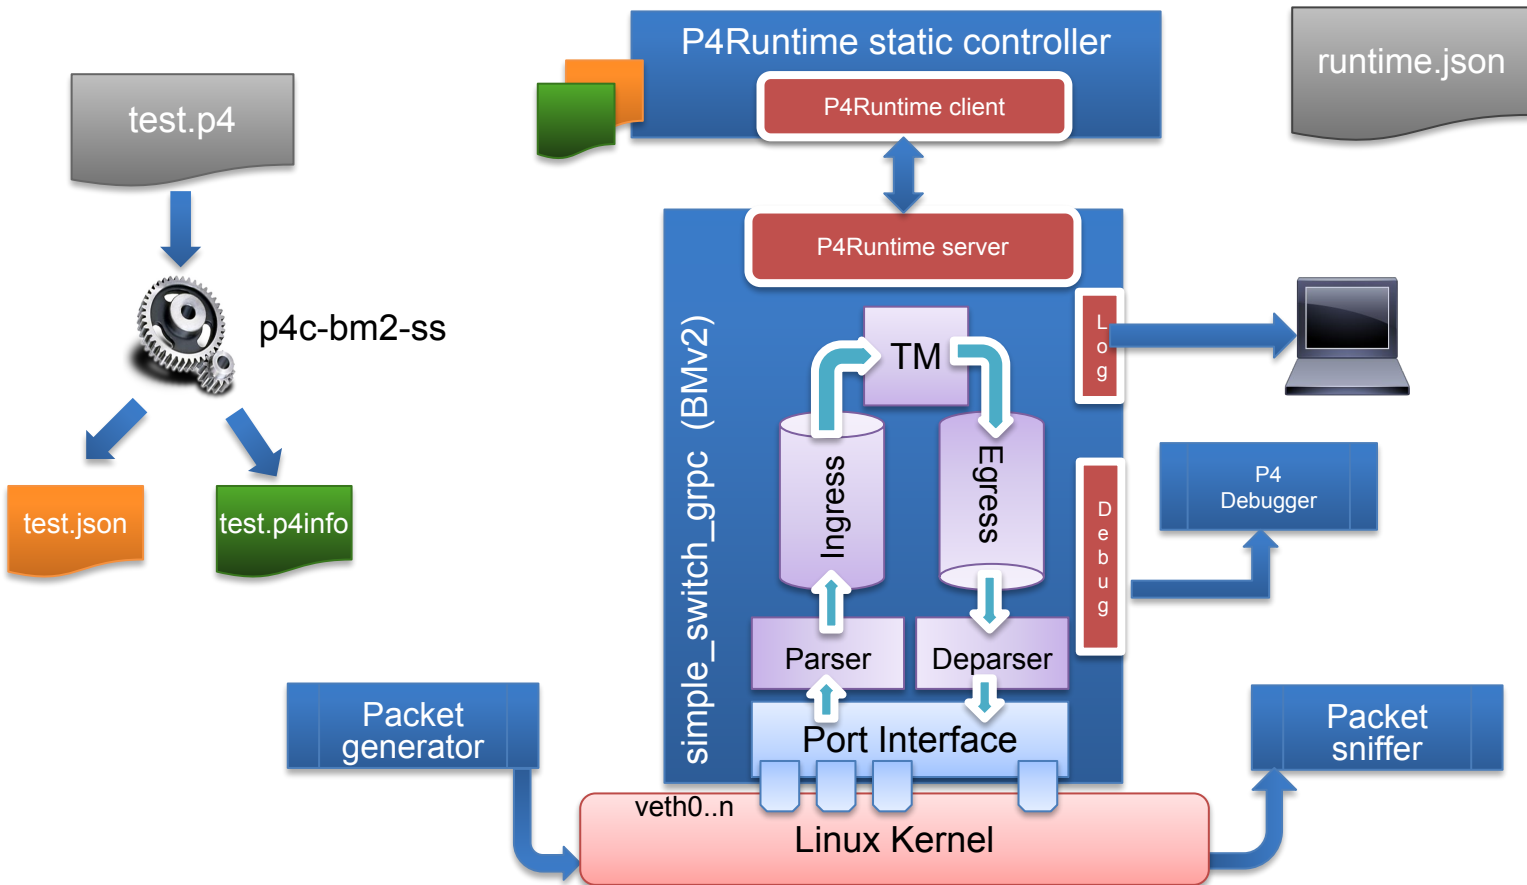 `* std_meta.egress_spec : 2`

Lab 2: P4Runtime

P4 Software Tools

Makefile: under the hood

test.p4

p4c-bm2-ss

test.json   test.p4info

P4Runtime static controller

P4Runtime client

runtime.json

simple_switch_grpc (BMv2)

P4Runtime server

TM

Log

Ingress   Egress

Debug

P4 Debugger

Parser   Deparser

Port Interface

Packet generator

veth0..n   Linux Kernel

Packet sniffer

Step 1: P4 Program compilation



```
$ p4c-bm2-ss --p4v 16 \
    -o test.json \
    --p4runtime-file test.p4info \
    --p4runtime-format text \
    test.p4
```

Step 2: Preparing veth Interfaces


```
$ sudo ~/p4lang/tutorials/examples/veth_setup.sh

# ip link add name veth0 type veth peer name veth1
# for iface in "veth0 veth1"; do
    ip link set dev ${iface} up
    sysctl net.ipv6.conf.${iface}.disable_ipv6=1
    TOE_OPTIONS="rx tx sg tso ufo gso gro lro rxvlan txvlan rxhash"
    for TOE_OPTION in $TOE_OPTIONS; do
        /sbin/ethtool --offload $intf "$TOE_OPTION"
    done
  done
```
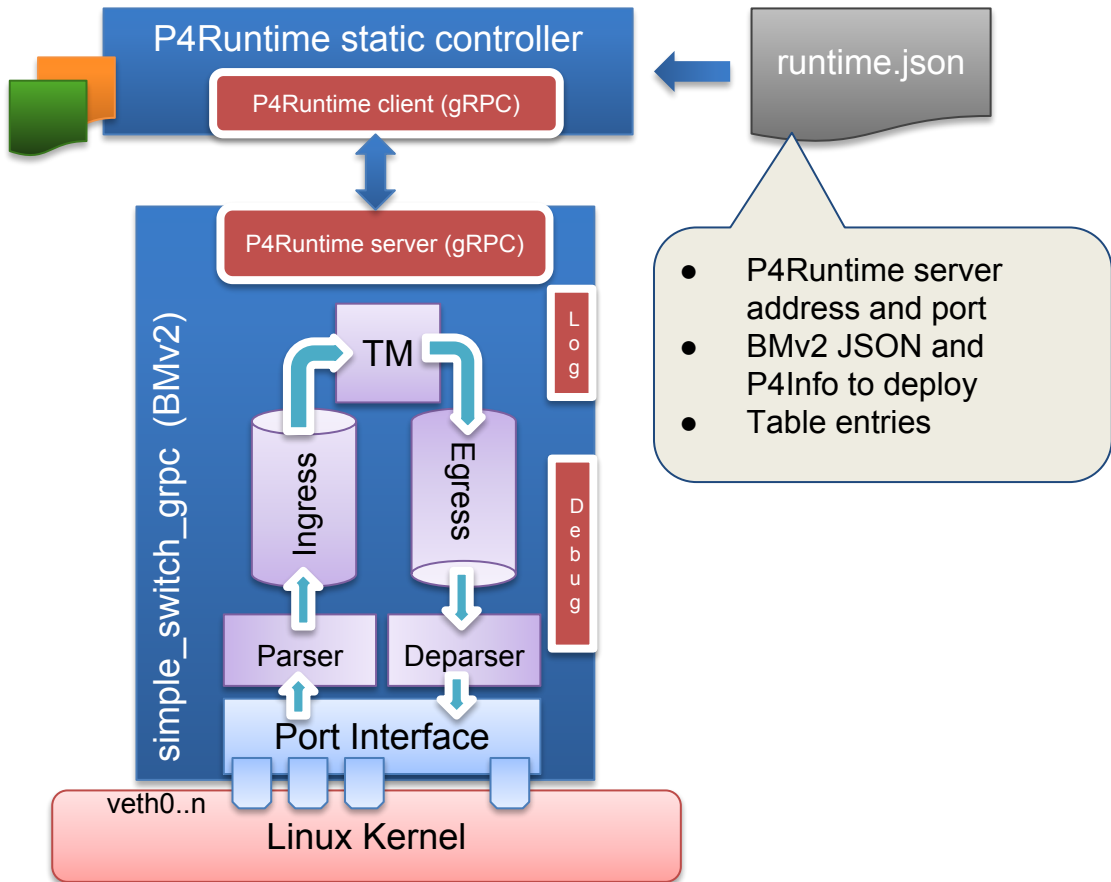
Step 3: Starting BMv2

```
$ sudo simple_switch_grpc --log-console --dump-packet-data 64 \
         -i 0@veth0 -i 1@veth2 … [--pcap]  --no-p4           \
         -- --grpc-server-addr 0.0.0.0:50051 --cpu-port 255   \
         test.json
```
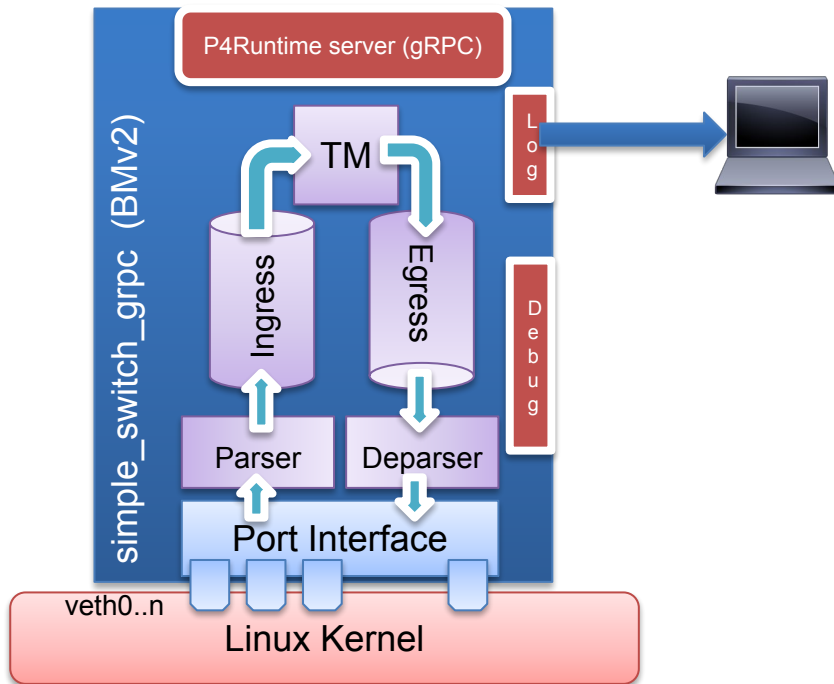
Step 4: Starting P4Runtime static controller


```
$ python $(RUN_SCRIPT) -t
$(TOPO) $(run_args)
```

- P4Runtime server address and port
- BMv2 JSON and P4Info to deploy
- Table entries

7

P4Runtime

- **API overview**
- **Workflow**
- **Exercise - Tunneling**

Runtime control of P4 data planes

User supplied

Vendor supplied

P4 Program

P4 Compiler

P4 Architecture Model

Target-specific configuration binary

Load

Control Plane

Add/remove table entries

Extern control

Packet-in/out

CPU port

Tables

Extern objects

Data Plane

Focus of this session
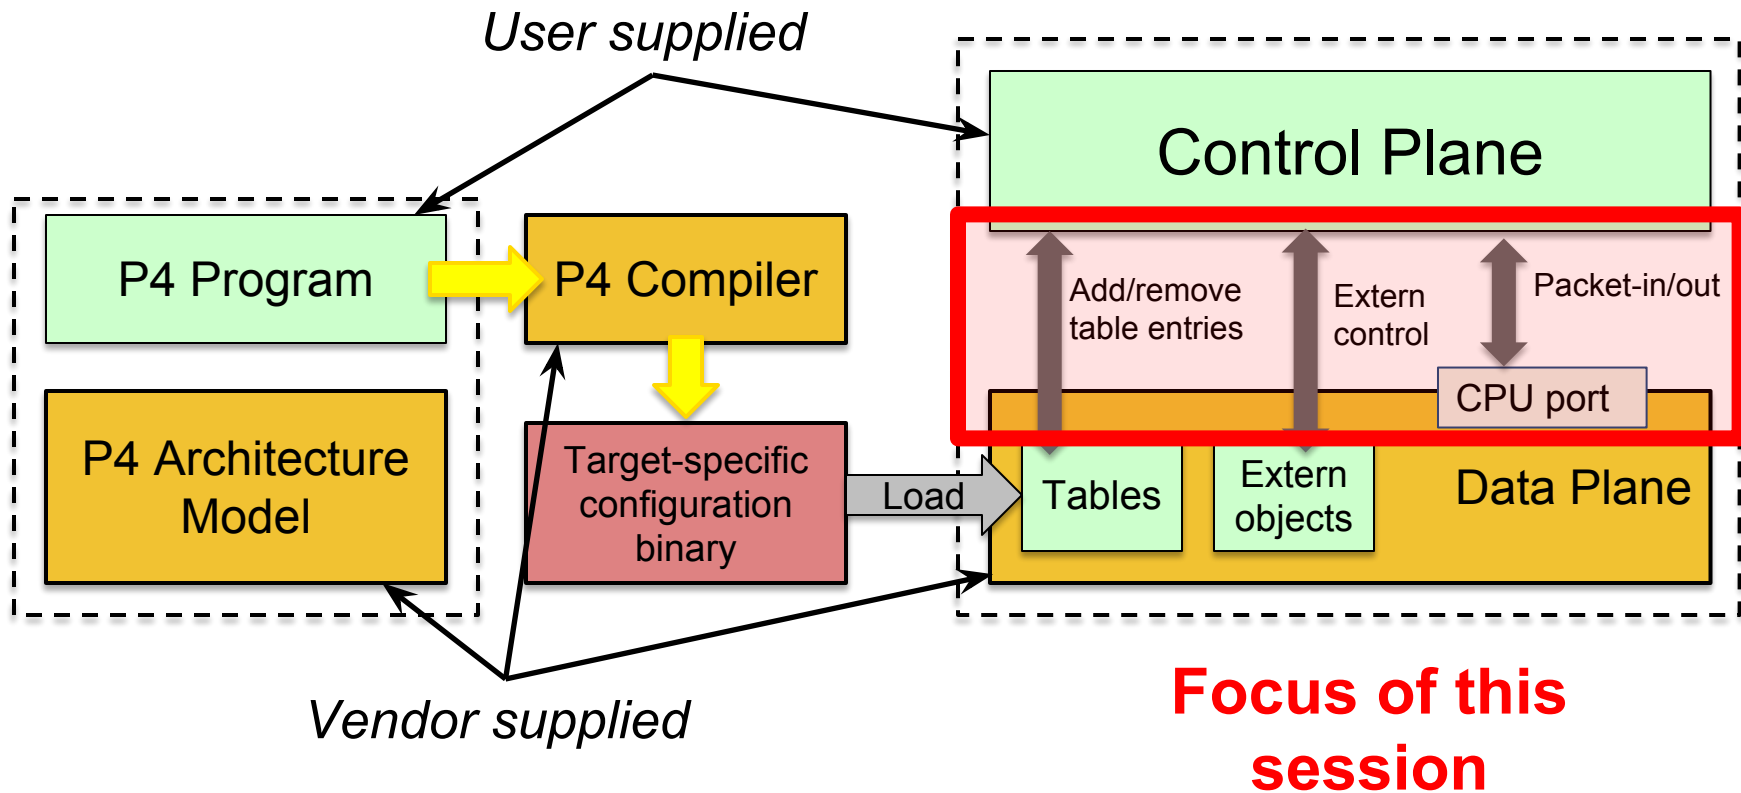
Existing approaches to runtime control

- **P4 compiler auto-generated runtime APIs**
  - Program-dependent -- hard to provision new P4 program without restarting the control plane!

- **BMv2 CLI**
  - Program-independent, but target-specific -- control plane not portable!

- **OpenFlow**
  - Target-independent, but protocol-dependent -- protocol headers and actions baked in the specification!

- **OCP Switch Abstraction Interface (SAI)**
  - Target-independent, but protocol-dependent

Why do we need another data plane control API?

Properties of a runtime control API

API	Target-independent	Protocol-independent
P4 compiler auto-generated	✔	✘
BMv2 CLI	✘	✔
OpenFlow	✔	✘
SAI	✔	✘
P4Runtime	✔	✔

What is P4Runtime?

- **Framework for runtime control of P4 targets**
  - Open-source API + server implementation
    https://github.com/p4lang/PI
  - Initial contribution by Google and Barefoot

- **Work-in-progress by the p4.org API WG**
  - Draft of version 1.0 available

- **Protobuf-based API definition**
  - p4runtime.proto
  - gRPC transport

- **P4 program-independent**
  - API doesn't change with the P4 program

- **Enables field-reconfigurability**
  - Ability to push new P4 program without recompiling the
    software stack of target switches

Control plane

p4runtime.proto
(API)

Program-independent
server (e.g. gRPC)

Target driver

P4 target

13
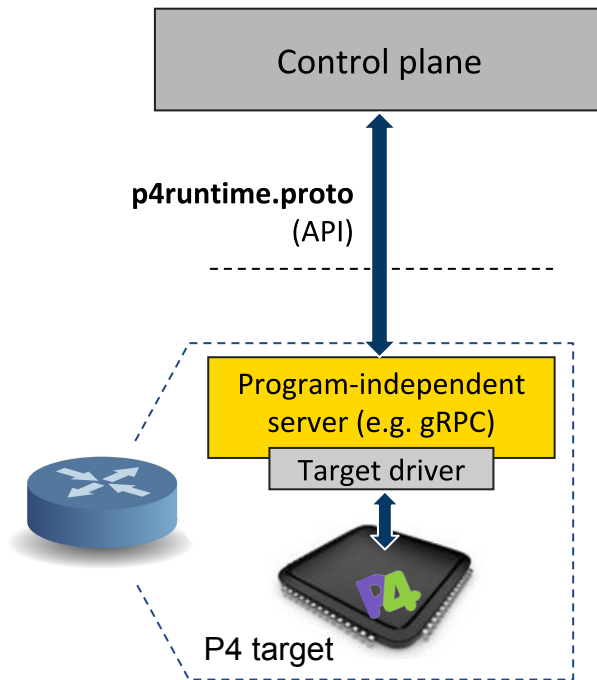
Protocol Buffers (protobuf) Basics

- Language for describing data for serialization in a structured way

- Common binary wire-format

- Language-neutral
  - Code generators for: *Action Script, C, C++, C#, Clojure, Lisp, D, Dart, Erlang, Go, Haskell, Java, Javascript, Lua, Objective C, OCaml, Perl, PHP, Python, Ruby, Rust, Scala, Swift, Visual Basic, ...*

- Platform-neutral

- Extensible and backwards compatible

- Strongly typed

```
syntax = "proto3";

message Person {
  string name = 1;
  int32 id = 2;
  string email = 3;

  enum PhoneType {
    MOBILE = 0;
    HOME = 1;
    WORK = 2;
  }

  message PhoneNumber {
    string number = 1;
    PhoneType type = 2;
  }

  repeated PhoneNumber phone = 4;
}
```

gRPC Basics

- Use Protocol Buffers to define service API and messages

- Automatically generate native stubs in:
  - C / C++
  - C#
  - Dart
  - Go
  - Java
  - Node.js
  - PHP
  - Python
  - Ruby

- Transport over HTTP/2.0 and TLS
  - Efficient single TCP connection implementation that supports bidirectional streaming

gRPC Service Example

```
// The greeter service definition.
service Greeter {
  // Sends a greeting
  rpc SayHello (HelloRequest) returns (HelloReply) {}
}

// The request message containing the user's name.
message HelloRequest {
  string name = 1;
}

// The response message containing the greetings
message HelloReply {
  string message = 1;
}
```

More details here: https://grpc.io/docs/guides/

P4Runtime Service

Enables a local or remote entity to arbitrate mastership, load the pipeline/program, send/receive packets, and read and write forwarding table entries, counters, and other P4 entities.

```
service P4Runtime {
  rpc Write(WriteRequest) returns (WriteResponse) {}
  rpc Read(ReadRequest) returns (stream ReadResponse) {}
  rpc SetForwardingPipelineConfig(SetForwardingPipelineConfigRequest)
      returns (SetForwardingPipelineConfigResponse) {}
  rpc GetForwardingPipelineConfig(GetForwardingPipelineConfigRequest)
      returns (GetForwardingPipelineConfigResponse) {}
  rpc StreamChannel(stream StreamMessageRequest)
      returns (stream StreamMessageResponse) {}
}
```

P4Runtime Service

P4Runtime Protobuf Definition:

https://github.com/p4lang/p4runtime/blob/master/proto/p4/v1/p4runtime.proto

Service Specification:

Working draft of version 1.0 is available now

https://p4.org/p4-spec/docs/P4Runtime-v1.0.0.pdf

P4Runtime `Write Request`

```
message WriteRequest {
 uint64 device_id = 1;
 uint64 role_id = 2;
 Uint128 election_id = 3;
 repeated Update updates = 4;
}
```

```
message Update {
 enum Type {
   UNSPECIFIED = 0;
   INSERT = 1;
   MODIFY = 2;
   DELETE = 3;
 }
 Type type = 1;
 Entity entity = 2;
}
```

```
message Entity {
 oneof entity {
   ExternEntry extern_entry = 1;
   TableEntry table_entry = 2;
   ActionProfileMember
         action_profile_member = 3;
   ActionProfileGroup
         action_profile_group = 4;
   MeterEntry meter_entry = 5;
   DirectMeterEntry direct_meter_entry = 6;
   CounterEntry counter_entry = 7;
   DirectCounterEntry direct_counter_entry = 8;
   PacketReplicationEngineEntry
         packet_replication_engine_entry = 9;
   ValueSetEntry value_set_entry = 10;
   RegisterEntry register_entry = 11;
 }
}
```

P4Runtime Table Entry

p4runtime.proto simplified excerpts:

```
message TableEntry {
  uint32 table_id;
  repeated FieldMatch match;
  Action action;
  int32 priority;
  ...
}
```

```
message Action {
  uint32 action_id;
  message Param {
    uint32 param_id;
    bytes value;
  }
  repeated Param params;
}
```

```
message FieldMatch {
  uint32 field_id;
  message Exact {
    bytes value;
  }
  message Ternary {
    bytes value;
    bytes mask;
  }
  ...
  oneof field_match_type {
    Exact exact;
    Ternary ternary;
    ...
  }
}
```

Full protobuf definition:
https://github.com/p4lang/PI/blob/master/proto/p4/p4runtime.proto

To add a table entry, the control plane needs to know:

- **IDs of P4 entities**
  - Tables, field matches, actions, params, etc.

- **Field matches for the particular table**
  - Match type, bitwidth, etc.

- **Parameters for the particular action**
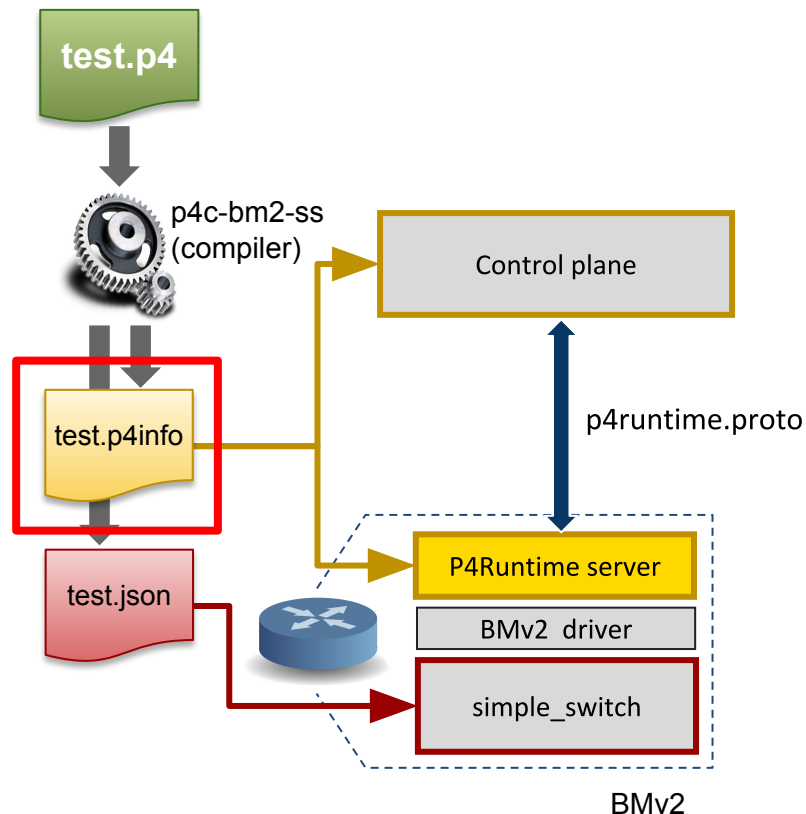
- **Other P4 program attributes**

P4Runtime workflow

P4Info

- **Captures P4 program attributes needed at runtime**
    - IDs for tables, actions, params, etc.
    - Table structure, action parameters, etc.

- **Protobuf-based format**

- **Target-independent compiler output**
    - Same P4Info for BMv2, ASIC, etc.

Full P4Info protobuf specification:
https://github.com/p4lang/p4runtime/blob/master/proto/p4/config/v1/p4info.proto

P4Info example

basic_router.p4

```
...

action ipv4_forward(bit<48> dstAddr,
                    bit<9> port) {
    /* Action implementation */
}

...

table ipv4_lpm {
    key = {
        hdr.ipv4.dstAddr: lpm;
    }
    actions = {
        ipv4_forward;
        ...
    }
    ...
}
```

P4 compiler

basic_router.p4info

```
actions {
  id: 16786453
  name: "ipv4_forward"
  params {
    id: 1
    name: "dstAddr"
    bitwidth: 48

    ...
    id: 2
    name: "port"
    bitwidth: 9
  }
}
...
tables {
  id: 33581985
  name: "ipv4_lpm"
  match_fields {
    id: 1
    name: "hdr.ipv4.dstAddr"
    bitwidth: 32
    match_type: LPM
  }
  action_ref_id: 16786453
}
```

P4Runtime Table Entry Example

basic_router.p4

```
action ipv4_forward(bit<48> dstAddr,
                    bit<9>  port) {
    /* Action implementation */
}
table ipv4_lpm {
    key = {
        hdr.ipv4.dstAddr: lpm;
    }
    actions = {
        ipv4_forward;
        ...
    }
    ...
}
```

**Control plane
generates**

Logical view of table entry

```
hdr.ipv4.dstAddr=10.0.1.1/32
      -> ipv4_forward(00:00:00:00:00:10, 7)
```

**Protobuf message**

```
table_entry {
  table_id: 33581985
  match {
    field_id: 1
    lpm {
      value: "\n\000\001\001"
      prefix_len: 32
    }
  }
  action {
    action_id: 16786453
    params {
      param_id: 1
      value: "\000\000\000\000\000\n"
    }
    params {
      param_id: 2
      value: "\000\007"
    }
  }
}
```

P4Runtime `SetPipelineConfig`

```
message SetForwardingPipelineConfigRequest {
  enum Action {
    UNSPECIFIED = 0;
    VERIFY = 1;
    VERIFY_AND_SAVE = 2;
    VERIFY_AND_COMMIT = 3;
    COMMIT = 4;
    RECONCILE_AND_COMMIT = 5;
  }
  uint64 device_id = 1;
  uint64 role_id = 2;
  Uint128 election_id = 3;
  Action action = 4;
  ForwardingPipelineConfig config = 5;
}
```
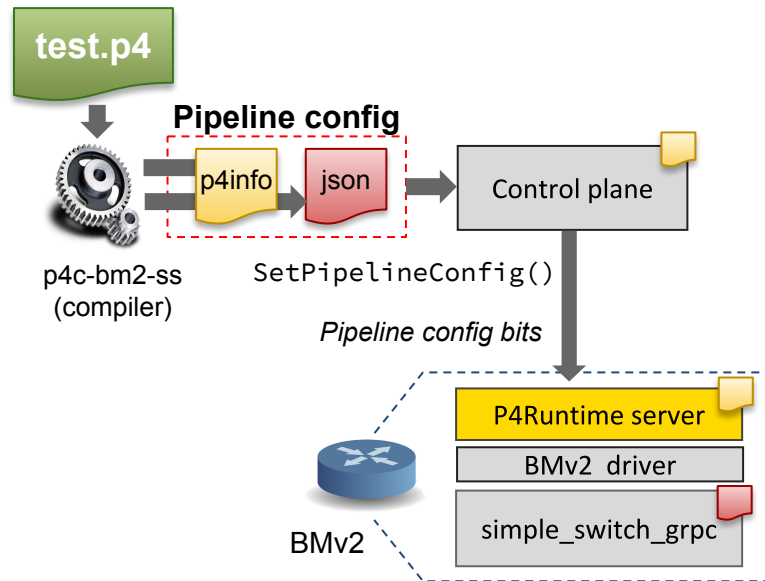


```
message ForwardingPipelineConfig {
  config.P4Info p4info = 1;
  // Target-specific P4 configuration.
  bytes p4_device_config = 2;
}
```

P4Runtime `StreamChannel`

```
message StreamMessageRequest {
 oneof update {
   MasterArbitrationUpdate
          arbitration = 1;
   PacketOut packet = 2;
   DigestListAck digest_ack = 3;
 }
}
```

```
// Packet sent from the controller to the switch.
message PacketOut {
 bytes payload = 1;
 // This will be based on P4 header annotated as
 // @controller_header("packet_out").
 // At most one P4 header can have this annotation.
 repeated PacketMetadata metadata = 2;
}
```

```
message StreamMessageResponse {
 oneof update {
   MasterArbitrationUpdate
          arbitration = 1;
   PacketIn packet = 2;
   DigestList digest = 3;
 }
}
```

```
// Packet sent from the switch to the controller.
message PacketIn {
 bytes payload = 1;
 // This will be based on P4 header annotated as
 // @controller_header("packet_in").
 // At most one P4 header can have this annotation.
 repeated PacketMetadata metadata = 2;
}
```

P4Runtime Common Parameters

- **`device_id`**
  - Specifies the specific forwarding chip or software bridge
  - **Set to 0 for single chip platforms**
- **`role_id`**
  - Corresponds to a role with specific capabilities (i.e. what operations, P4 entities, behaviors, etc. are in the scope of a given role)
  - Role definition is currently agreed upon between control and data planes offline
  - **Default role_id (0) has full pipeline access**
- **`election_id`**
  - P4Runtime supports mastership on a per-role basis
  - Client with the highest election ID is referred to as the "master", while all other clients are referred to as "slaves"
  - **Set to 0 for single instance controllers**

Mastership Arbitration

- **Upon connecting to the device, the client (e.g. controller) needs to open a `StreamChannel`**
- **The client must advertise its `role_id` and `election_id` using a `MasterArbitrationUpdate` message**
  - If `role_id` is not set, it implies the default role and will be granted full pipeline access
  - The `election_id` is opaque to the server and determined by the control plane (can be omitted for single-instance control plane)
- **The switch marks the client for each role with the highest `election_id` as master**
- **Master can:**
  - Perform `Write` requests
  - Receive `PacketIn` messages
  - Send `PacketOut` messages

Remote control
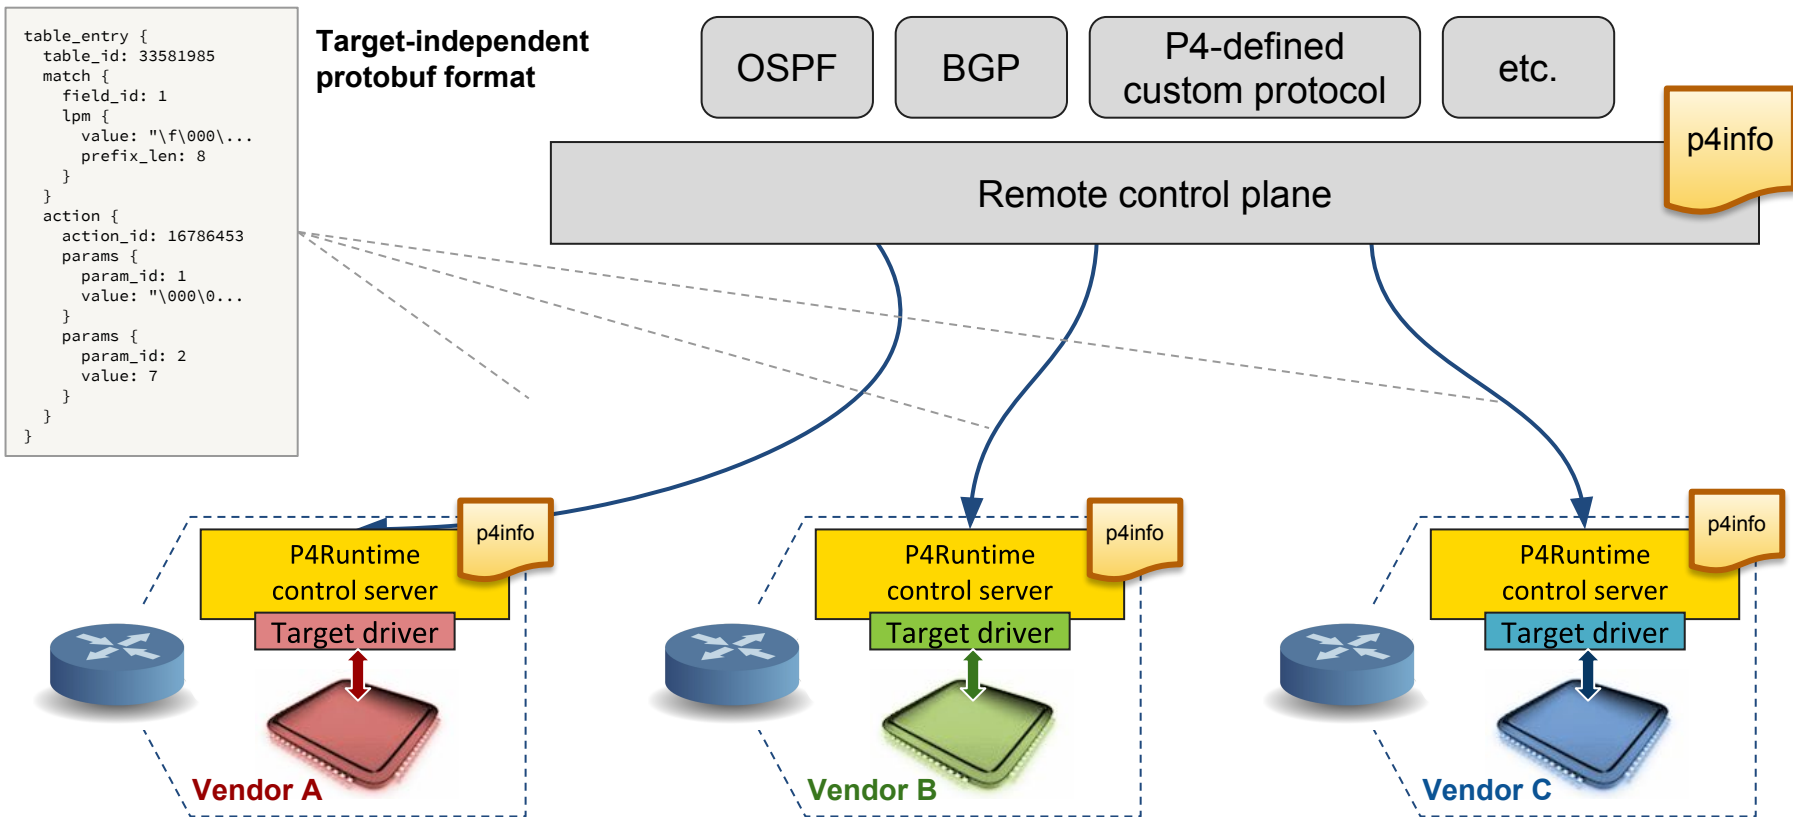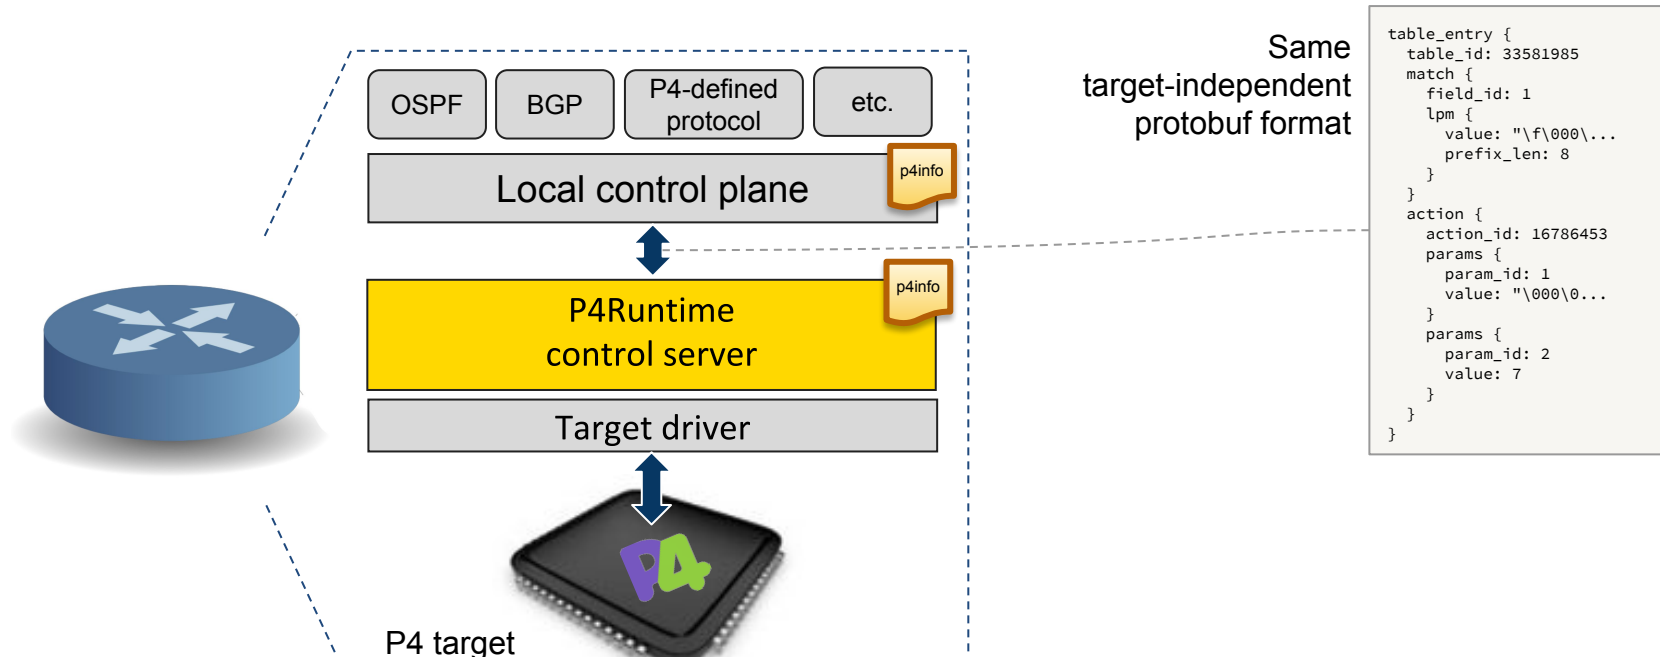
```
table_entry {
  table_id: 33581985
  match {
    field_id: 1
    lpm {
      value: "\f\000\...
      prefix_len: 8
    }
  }
  action {
    action_id: 16786453
    params {
      param_id: 1
      value: "\000\0...
    }
    params {
      param_id: 2
      value: 7
    }
  }
}
```

Target-independent protobuf format

| OSPF | BGP | P4-defined custom protocol | etc. |

Remote control plane

p4info

P4Runtime control server

p4info

Target driver

Vendor A

P4Runtime control server

p4info

Target driver

Vendor B

P4Runtime control server

p4info

Target driver

Vendor C

Local control

P4Runtime on Fixed-Function Chips

TODO pull some stratum slides

P4 Program as Fixed-Function Chip Abstraction

- P4 program tailored to apps / role - does not describe the hardware
- Switch maps program to fixed-function ASIC
- Enables portability

P4Runtime API recap

Things we covered:

- **P4Info**
- **Table entries**
- **Set pipeline config**

What we didn't cover:

- **How to control other P4 entities**
  - Externs, counters, meters
- **Packet-in/out support**
- **Controller replication**
  - Via master-slave arbitration
- **Batched reads/writes**
- **Switch configuration**
  - Outside the P4Runtime scope
  - Achieved with other mechanisms
    - e.g., OpenConfig and gNMI

Work-in-progress by the p4.org API WG
Expect API changes in the future

P4 compiler generates 2 files:

1. Target-specific binaries
   - Used to configure switch pipeline
     (e.g. binary config for ASIC, bitstream for FPGA, etc.)

2. P4Info file
   - Captures P4 program attributes needed to runtime control
     - Tables, actions, parameters, etc.
   - Protobuf-based format
   - Target-independent compiler output
     - Same P4Info for SW switch, ASIC, etc.

Full P4Info protobuf specification:
https://github.com/p4lang/PI/blob/master/proto/p4/config/p4info.proto

33
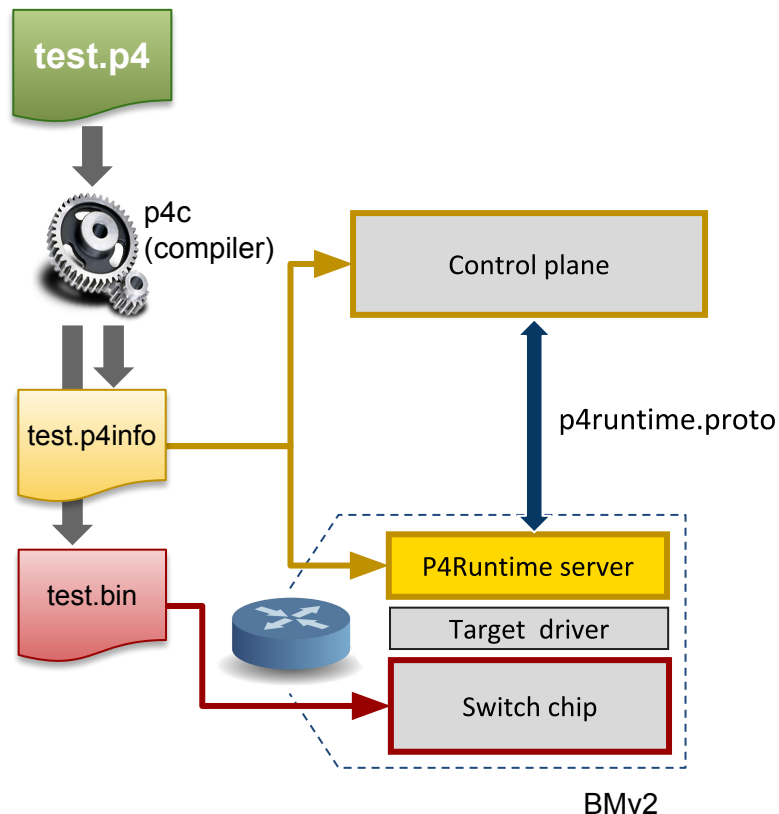
P4Runtime exercise

Exercise Overview

Controller's responsibilities:

1. **Establish a gRPC connection to the switches for the P4Runtime service**
2. **Push the P4 program to each switch**
3. **Write the tunnel forwarding rules:**
   a. `myTunnel_ingress` rule to encapsulate packets on the ingress switch
   b. `myTunnel_forward` rule to forward packets on the ingress switch
   c. `myTunnel_egress` rule to decapsulate and forward packets on the egress switch
4. **Read the tunnel ingress and egress counters every 2 seconds**

mycontroller.py

h1
(10.0.1.1)
(dst_id: 1)

s1

h2
(10.0.2.2)
(dst_id: 2)

s2

s3

h3
(10.0.3.3) (dst_id: 3)

Getting started

The source code has already been downloaded on your VM:
`~/tutorials/exercises/p4runtime`

You should start by reading the `README.md`

In this exercise, you will need to complete the implementation of `writeTunnelRules` in `mycontroller.py`

You will need two Terminal windows: one for your dataplane network (Mininet) that you will start using `make`, and the other is for your controller program.

To find the source code:
`https://github.com/p4lang/tutorials/`

📖 README.md

Implementing a Control Plane using P4 Runtime

Introduction

In this exercise, we will be using P4 Runtime to send flow entries to the switch instead of using the switch's CLI. We will be building on the same P4 program that you used in the basic_tunnel exercise. The P4 program has been renamed to `advanced_tunnel.py` and has been augmented with two counters (`ingressTunnelCounter` , `egressTunnelCounter`) and two new actions (`myTunnel_ingress` , `myTunnel_egress`).

You will use the starter program, `mycontroller.py` , and a few helper libraries in the `p4runtime_lib` directory to create the table entries necessary to tunnel traffic between host 1 and 2.

> **Spoiler alert:** There is a reference solution in the `solution` sub-directory. Feel free to compare your implementation to the reference.

Step 1: Run the (incomplete) starter code

The starter code for this assignment is in a file called `mycontroller.py` , and it will install only some of the rules that you need to tunnel traffic between two hosts.

Let's first compile the new P4 program, start the network, use `mycontroller.py` to install a few rules, and look at the `ingressTunnelCounter` to see that things are working as expected.

1. In your shell, run:

```
make
```

Example: ONOS

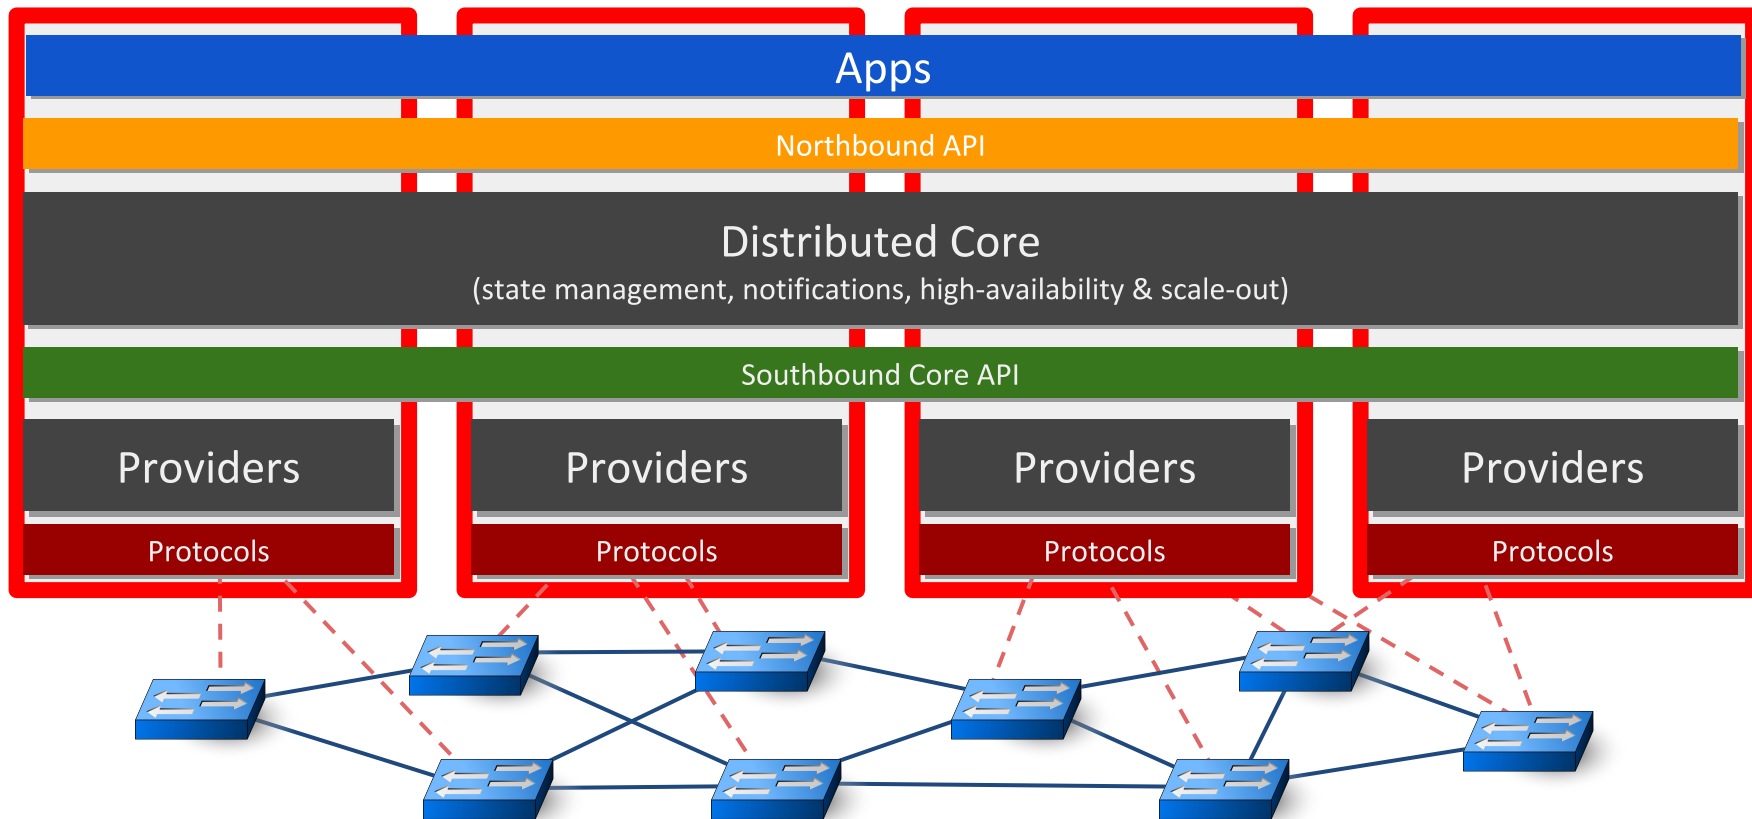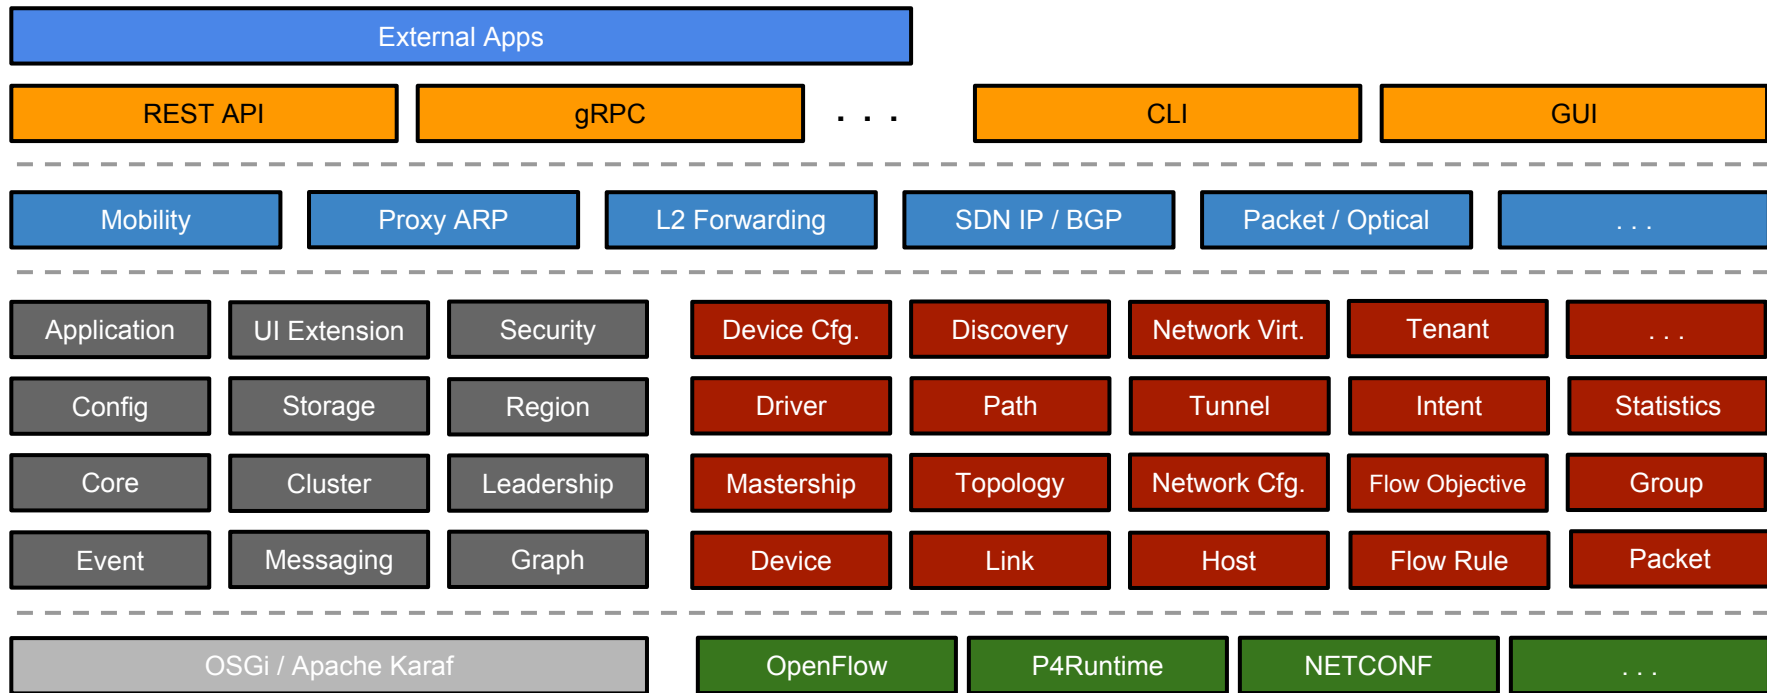*Open Network Operating System (ONOS) is an open source SDN network operating system, originally created by ON.Lab and currently hosted by the Linux Foundation.*
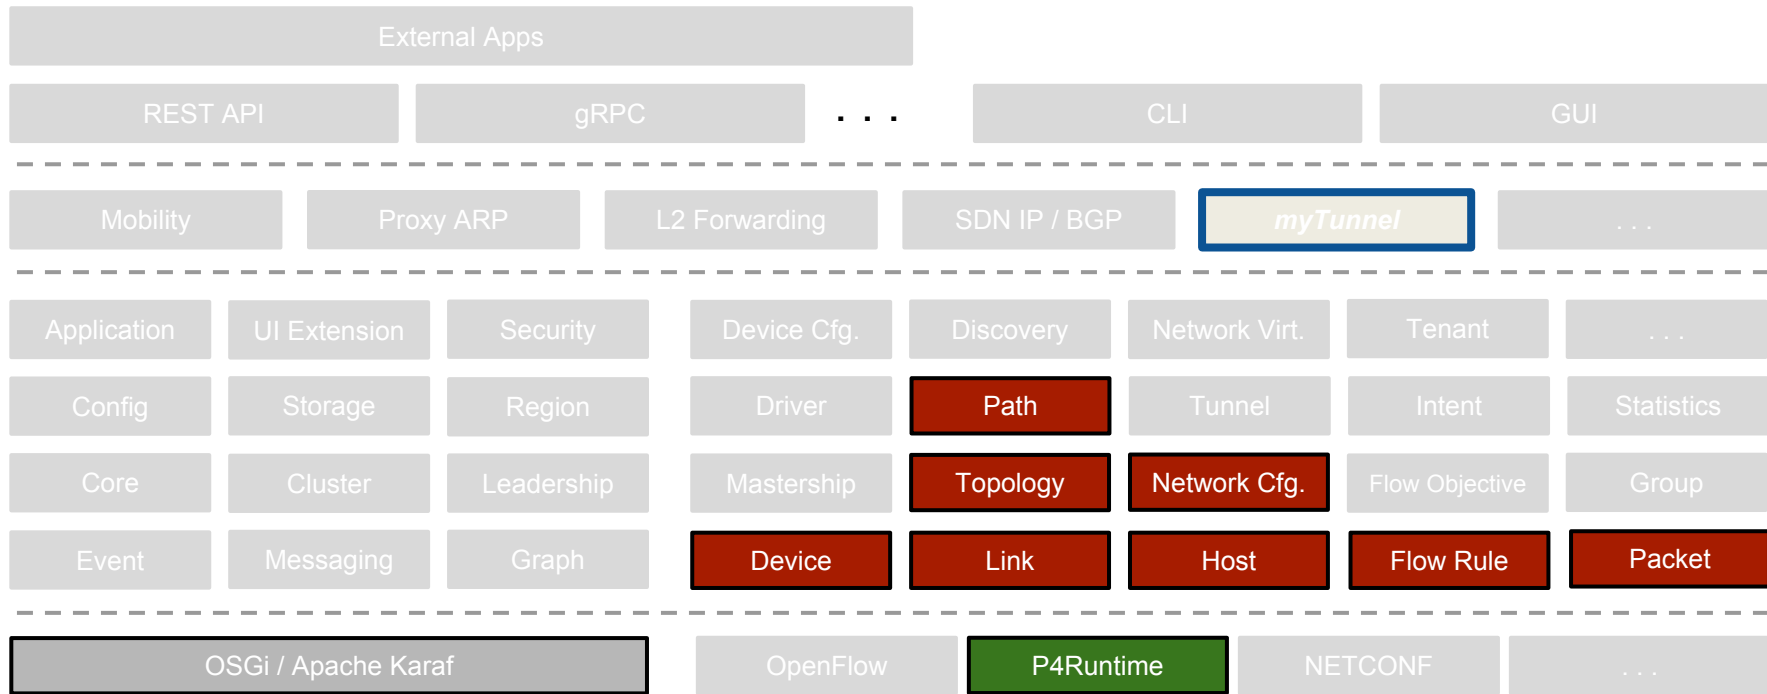
ONOS – SDN Controller

Apps

Northbound API

Distributed Core
(state management, notifications, high-availability & scale-out)

Southbound Core API

Providers

Providers

Providers

Providers

Protocols

Protocols

Protocols

Protocols

ONOS Core Subsystems

External Apps			

REST API	gRPC	. . .	CLI	GUI

Mobility	Proxy ARP	L2 Forwarding	SDN IP / BGP	Packet / Optical	. . .

Application	UI Extension	Security	Device Cfg.	Discovery	Network Virt.	Tenant	. . .
Config	Storage	Region	Driver	Path	Tunnel	Intent	Statistics
Core	Cluster	Leadership	Mastership	Topology	Network Cfg.	Flow Objective	Group
Event	Messaging	Graph	Device	Link	Host	Flow Rule	Packet

OSGi / Apache Karaf	OpenFlow	P4Runtime	NETCONF	. . .

ONOS Core Subsystems

External Apps				
REST API	gRPC	. . .	CLI	GUI

Mobility	Proxy ARP	L2 Forwarding	SDN IP / BGP	*myTunnel*	. . .

Application	UI Extension	Security	Device Cfg.	Discovery	Network Virt.	Tenant	. . .
Config	Storage	Region	Driver	**Path**	Tunnel	Intent	Statistics
Core	Cluster	Leadership	Mastership	**Topology**	**Network Cfg.**	Flow Objective	Group
Event	Messaging	Graph	**Device**	**Link**	**Host**	**Flow Rule**	**Packet**

| OSGi / Apache Karaf	OpenFlow	**P4Runtime**	NETCONF	. . .

Copyright © 2018 – P4.org, ONF

P4Runtime takeaways

- **Program-independent API**
  - API doesn't change with the P4 program
  - No need to restart the control-plane with a different P4 program

- **Device does not need to be fully programmable**
  - Can be used on fixed-function devices
  - Provided that their behavior can be expressed in P4

- **Protobuf-based format**
  - Well-supported serialization format in many languages
  - Supported by many RPC frameworks, e.g., gRPC
    - Auto-generate client/server code for different languages
    - No need to define common RPC features (e.g., authentication)

# P4 support in ONOS
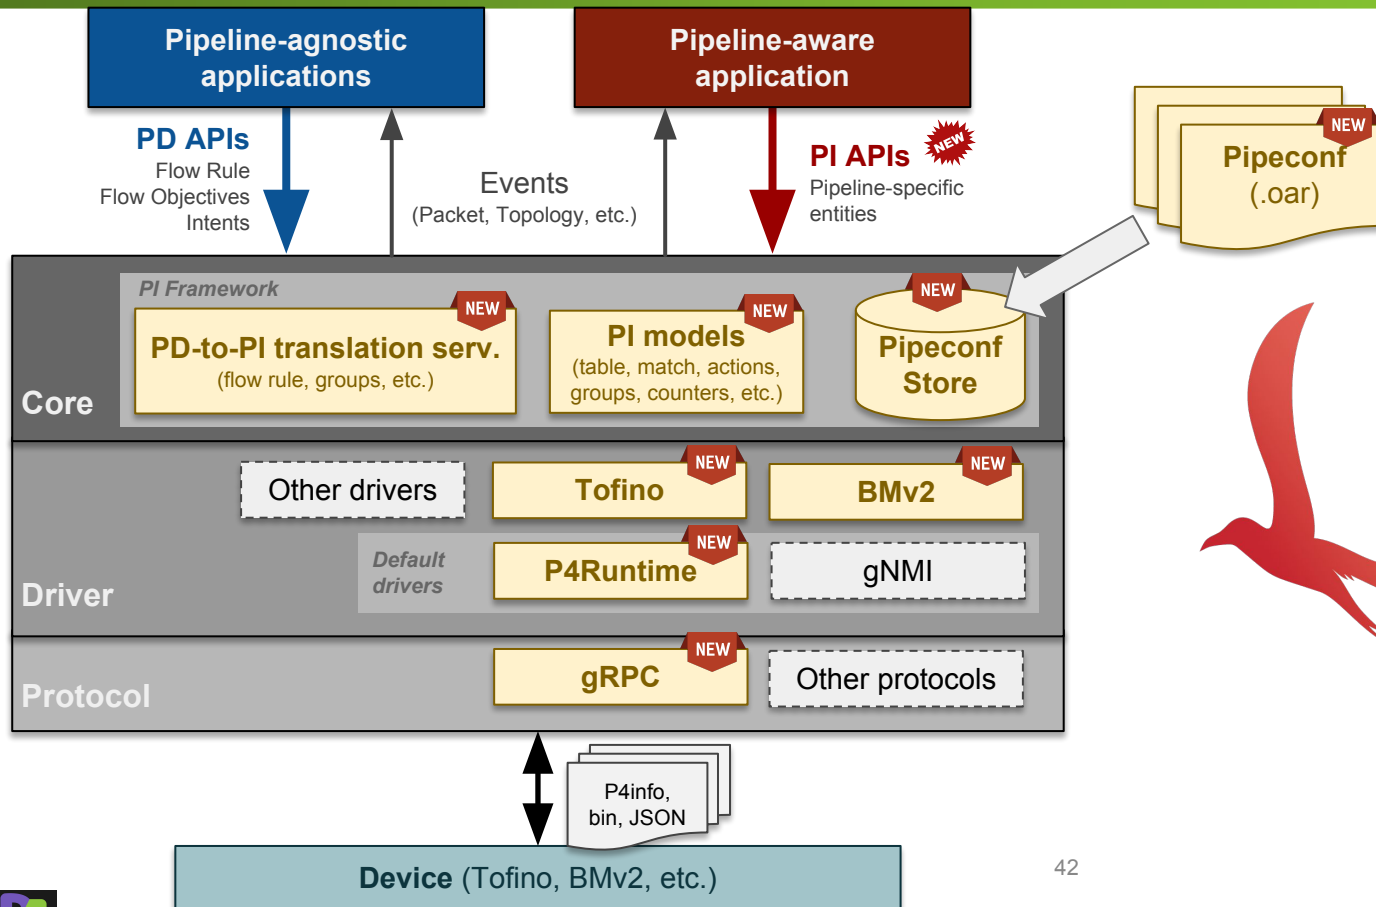
Lab 3: Advanced P4 Features

Plan

In this afternoon lab session, we'll explore some of P4's advanced features and see how to implement functionality commonly found in modern data planes.

- *Source Routing:* implement simple form of source routing using P4's header stacks
- *Load Balancing:* implement ECMP-like functionality using P4's `extern` functions
- *Network Telemetry:* implement fine-grained telemetry directly in the data plane

Source Routing

2,3,2,2,1, {payload}

Source Routing: Packet Format

```
#define MAX_HOPS 9

const bit<16> TYPE_IPV4 = 0x800;
const bit<16> TYPE_SRCROUTING = 0x1234;
header srcRoute_t {
  bit<1>    bos;
  bit<15>   port;
}

struct headers {
  ethernet_t          ethernet;
  srcRoute_t[MAX_HOPS] srcRoutes;
  ipv4_t              ipv4;
}
```

- **Parse source routes only if etherType is 0x1234**

- **The special value bos == 1 indicates the "bottom of stack"**

- **Forward packets using source routes, and also decrement IPv4 TTL**

- **Drop the packet if source routes are not valid**

- **Hint: Use the next, pop_front primitives packet.extract(hdr.srcRoutes.next) hdr.srcRoutes.pop_front(1)**

Simple Load Balancing

h1 → 10.0.0.1

h1
(10.0.1.1)

s1

s2

h2
(10.0.2.2)

s3

h3
(10.0.3.3)

Virtual IP (10.0.0.1)
shared by h2 and h3

Hashing (V1Model)

```
enum HashAlgorithm {
  csum16,
  xor16,
  crc32,
  crc32_custom,
  crc16,
  crc16_custom,
  random,
  identity
}
extern void hash<O, T, D, M>(
  out O result,
  in HashAlgorithm algo,
  in T base,
  in D data,
  in M max);
```

**Computes the hash of** data **(using** algo**) modulo** max **and adds it to** base

Uses type variables (like C++ templates / Java Generics) to allow hashing primitive to be used with many different types.
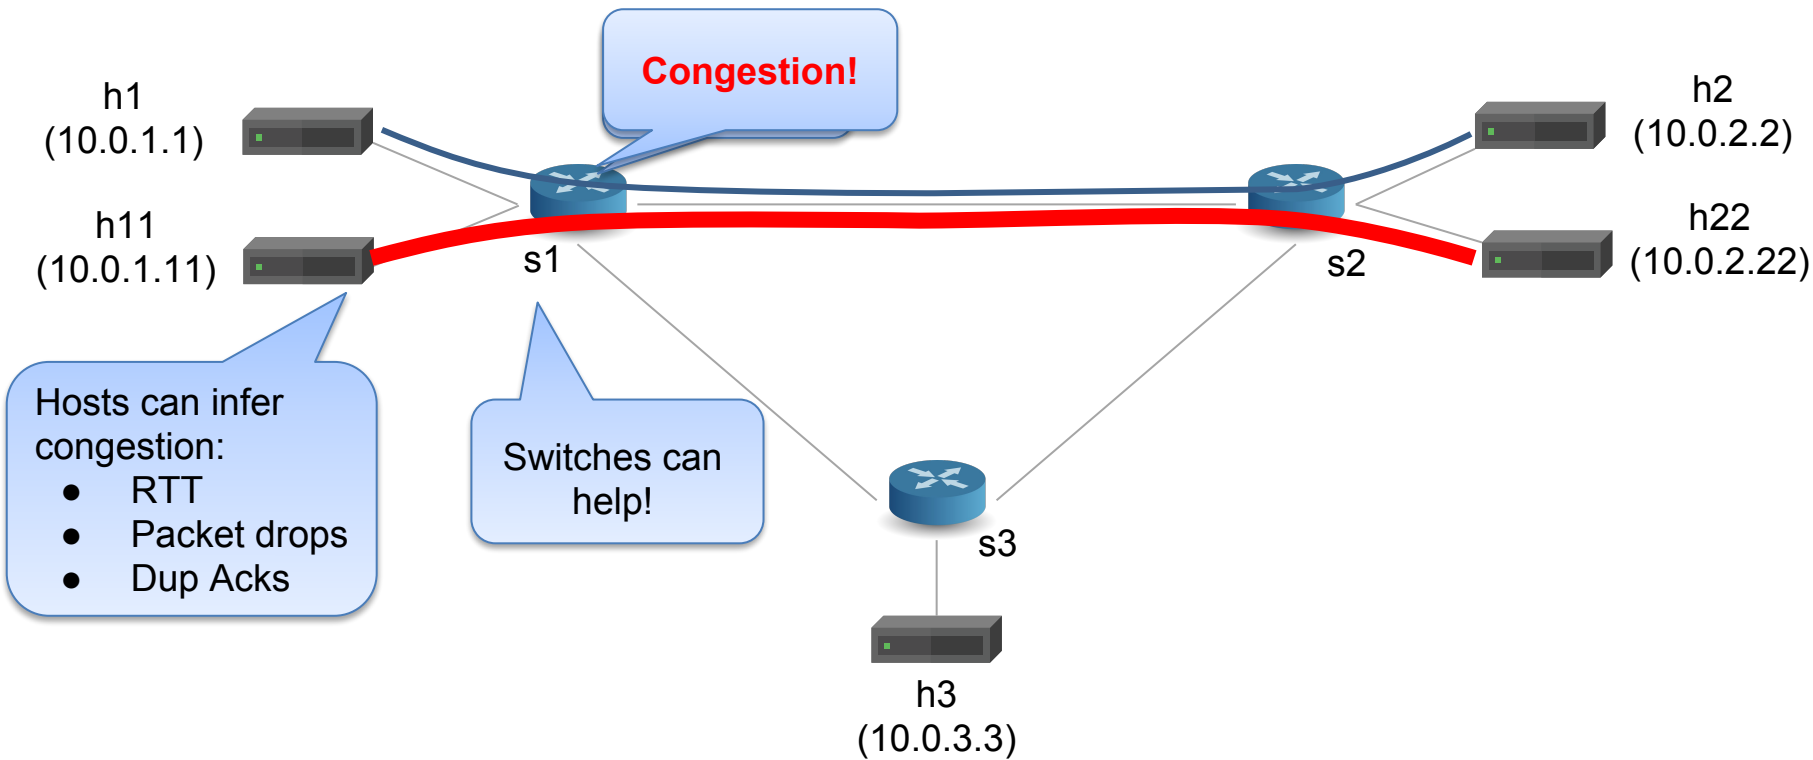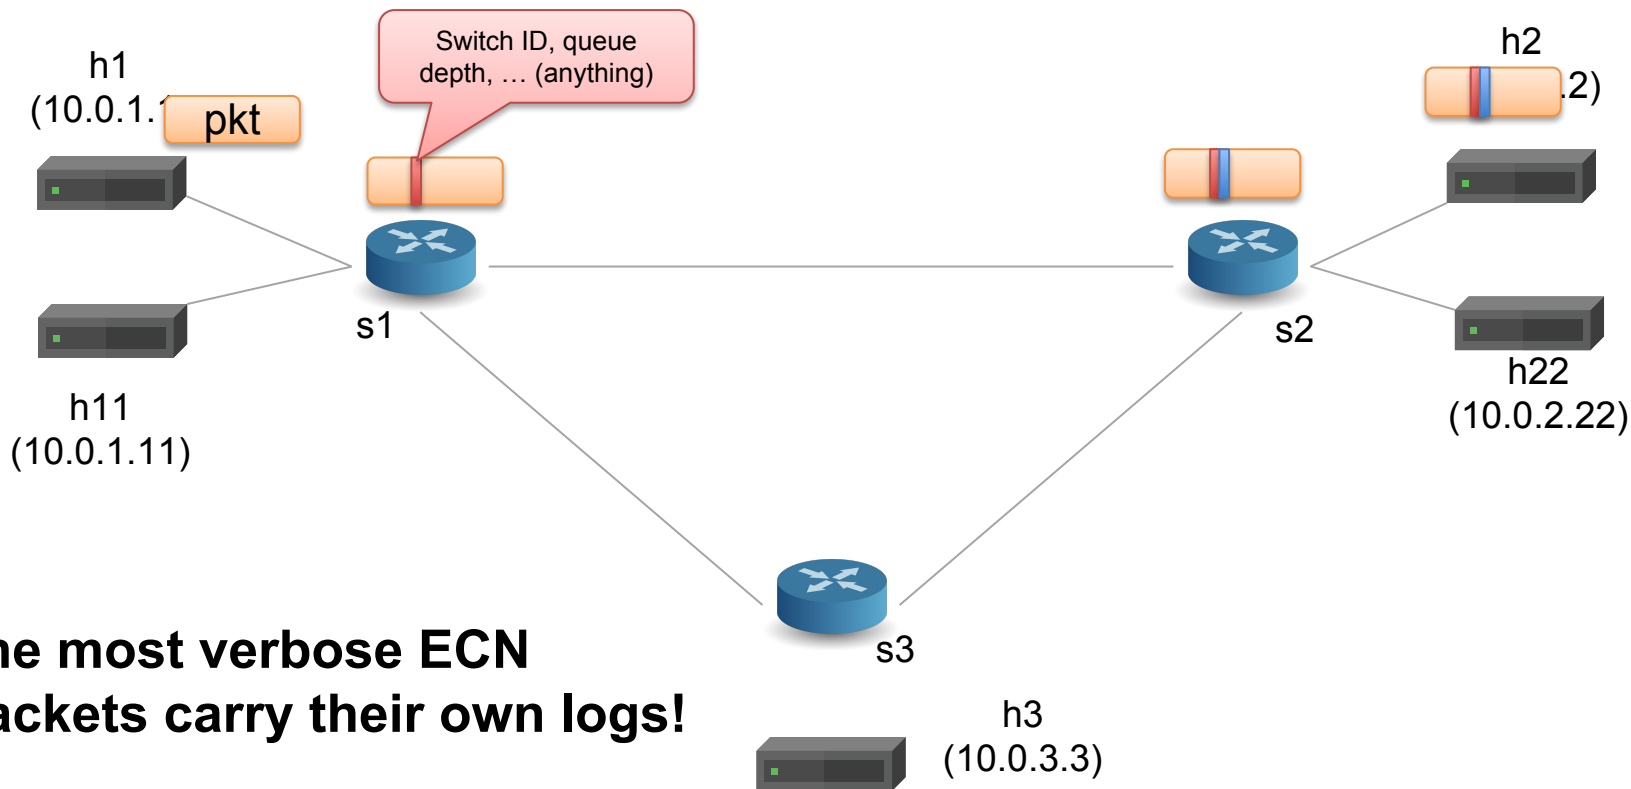
Monitoring & Debugging

Monitoring & Debugging

s1

Multi-Route Inspection

The most verbose ECN
Packets carry their own logs!

Multi-Route Inspect: Packet Format

```
header mri_t {
  bit<16>  count;
}
```

```
header switch_t {
  switchID_t  swid;
  qdepth_t    qdepth;
}
```

```
struct headers {
  ethernet_t        ethernet;
  ipv4_t            ipv4;
  ipv4_option_t     ipv4_option;
  mri_t             mri;
  switch_t[MAX_HOPS] swtraces;
}
```

- **Header validity opertions:**
  - hdr.setValid(): add_header
  - hdr.setInvalid(): remove_header
  - hdr.isValid():test validity

- **Header Stacks**
  - hdr[CNT] stk;
- **Header Stacks in Parsers**
  - stk.next
  - stk.last
  - stk.lastIndex
- **Header Stacks in Controls**
  - stk[i]
  - stk.size
  - stk.push_front(int count)
  - stk.pop_front(int count)

Header verification

```
/* Standard errors, defined in core.p4 */
error {
    NoError,              // no error
    PacketTooShort,       // not enough bits in packet for extract
    NoMatch,              // match expression has no matches
    StackOutOfBounds,     // reference to invalid element of a header stack
    OverwritingHeader,    // one header is extracted twice
    HeaderTooShort,       // extracting too many bits in a varbit field
    ParserTimeout         // parser execution time limit exceeded
}
/* Additional error added by the programmer */
error { IPv4BadHeader }
...
state parse_ipv4 {
  packet.extract(hdr.ipv4);
  verify(hdr.ipv4.version == 4, error.IPv4BadHeader);
  transition accept;
}
```

Wrapping up & Next Steps

Why P4$_{16}$?

- **Clearly defined semantics**
  - You can describe what your data plane program is doing
- **Expressive**
  - Supports a wide range of architectures through standard methodology
- **High-level, Target-independent**
  - Uses conventional constructs
  - Compiler manages the resources and deals with the hardware
- **Type-safe**
  - Enforces good software design practices and eliminates "stupid" bugs
- **Agility**
  - High-speed networking devices become as flexible as any software
- **Insight**
  - Freely mixing packet headers and intermediate results

Things we covered

- **The P4 "world view"**
  - Protocol-Independent Packet Processing
  - Language/Architecture Separation
  - If you can interface with it, it can be used
- **Key data types**
- **Constructs for packet parsing**
  - State machine-style programming
- **Constructs for packet processing**
  - Actions, tables and controls
- **Packet deparsing**
- **Architectures & Programs**

Things we didn't cover

- **Mechanisms for modularity**
  - Instantiating and invoking parsers or controls
- **Details of variable-length field processing**
  - Parsing and deparsing of options and TLVs
- **Architecture definition constructs**
  - How these "templated" definitions are created
- **Advanced features**
  - How to do learning, multicast, cloning, resubmitting
  - Header unions
- **Control plane interface**

The P4 Language Consortium

- **Consortium of academic and industry members**

- **Open source, evolving, domain-specific language**

- **Permissive Apache license, code on GitHub today**

- **Membership is free: contributions are welcome**

- **Independent, set up as a California nonprofit**

5